

Lecture 3 – Supervised Learning

COMPX310 | Irfan Ahmad | AI | Data Science | Computer Science



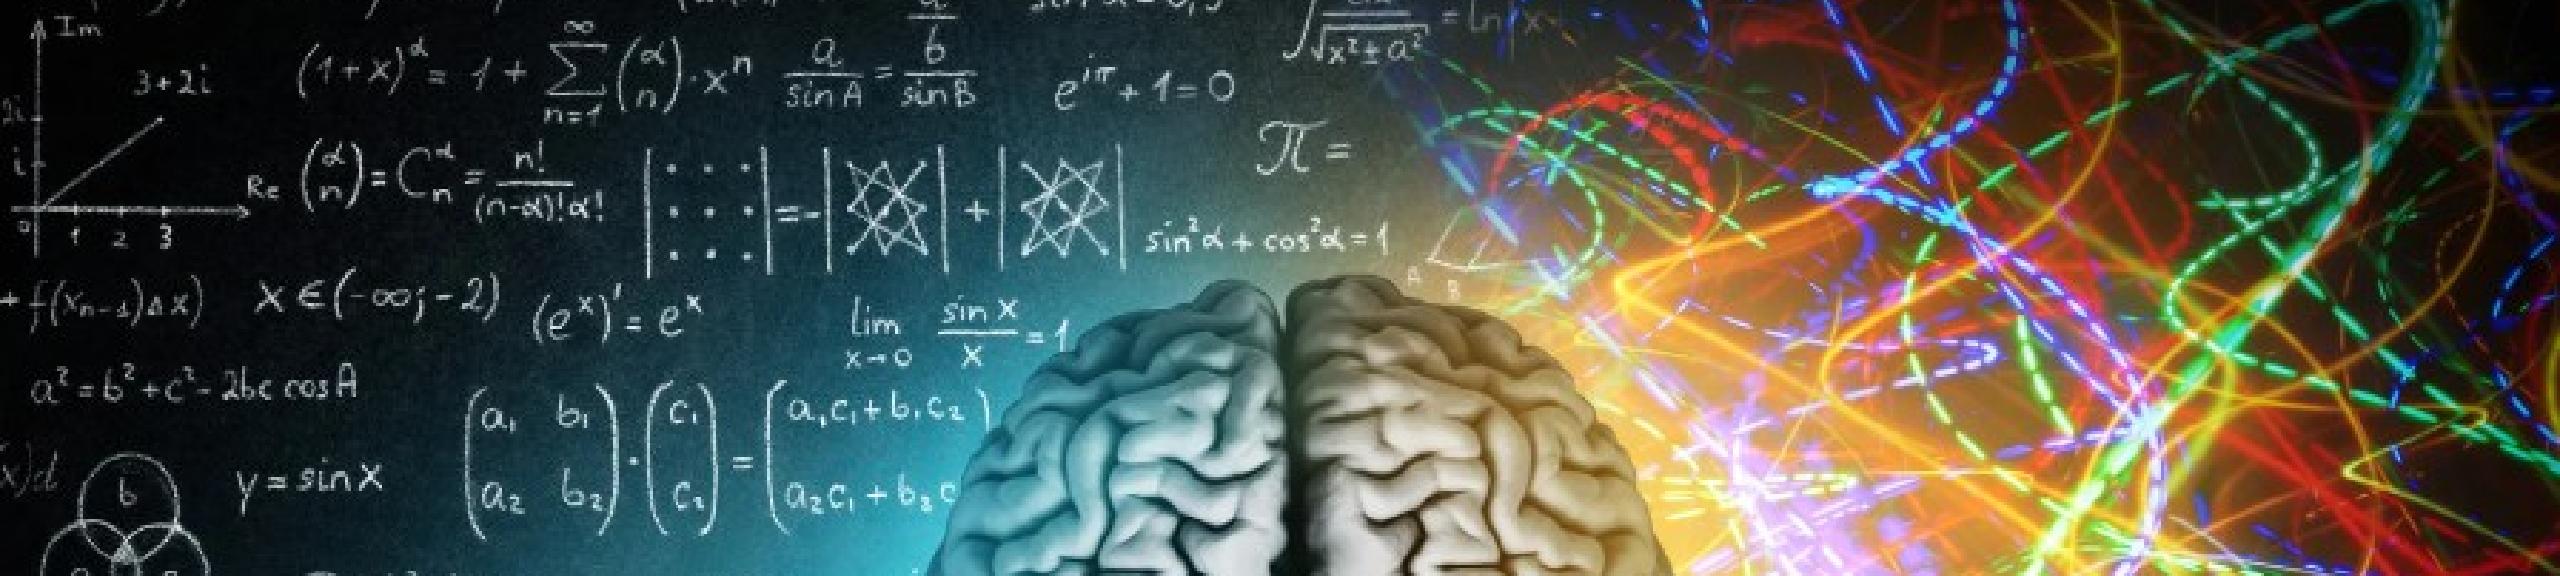
THE UNIVERSITY OF
WAIKATO
Te Whare Wānanga o Waikato

A collage of machine learning concepts:

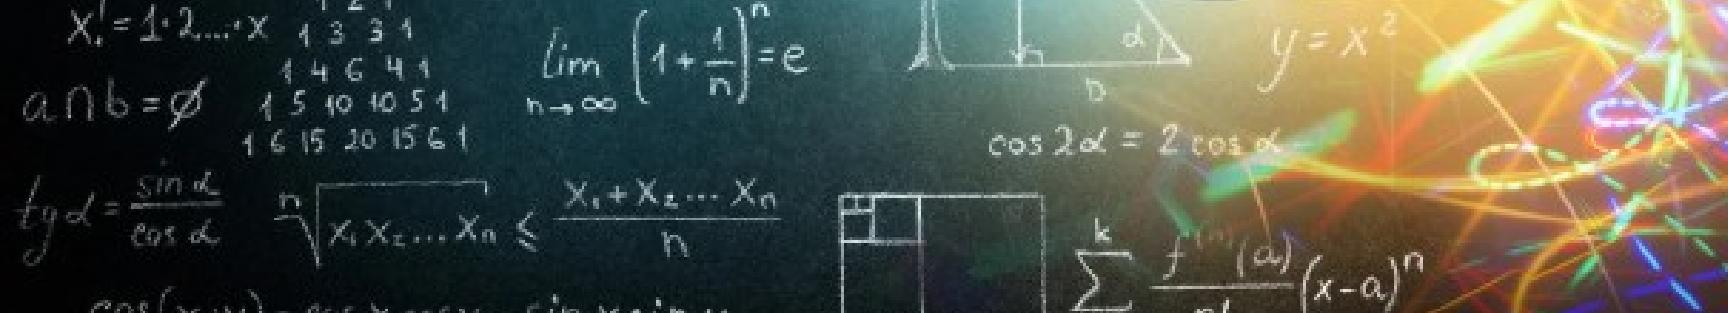
- Top Left: A scatter plot with red 'X' and blue hexagon data points, labeled $y = H|r(y)$.
- Top Center: A large glowing blue sign reading "MACHINE LEARNING".
- Top Right: Mathematical equations:
 - $P(x|y) = \frac{P(y|x)}{P(y)}$
 - $P(x|y) = \frac{(y)}{p(x)} \log, (x, x)$
 - $P(Y(x)|y, y) = \frac{a}{p(x|y)}$
 - $P(y|y) = p((x|y)) X(y)$
 - $P(x|y|x, y) = \frac{a}{\log(p|y)}$
 - $Y(x) = \sum_{y=1}^d P\left(\frac{x}{X} g(xy)\right)$
 - $P^+(x|y) = \frac{I(X - \sigma|+)}{P(x|y|y|y)}$
- Bottom Left: Three stylized gears in light blue, orange, and teal, each containing a geometric pattern (triangles, squares, circles).
- Bottom Center: Mathematical equations:
 - $f_x = \frac{\log, (p, y)}{\log (x y)}$
 - $x|y) = \frac{P(x|y)}{P(x|y)y|x|y)} = \frac{d}{P(x|x|y)}$
 - $x = f(t) = \frac{1}{x}$

Today's Topics

- Overview of Supervised Learning
- Classification Example 1 (MNIST Data)
- Classification Example 2 (KNN Method)



Overview of Supervised Learning



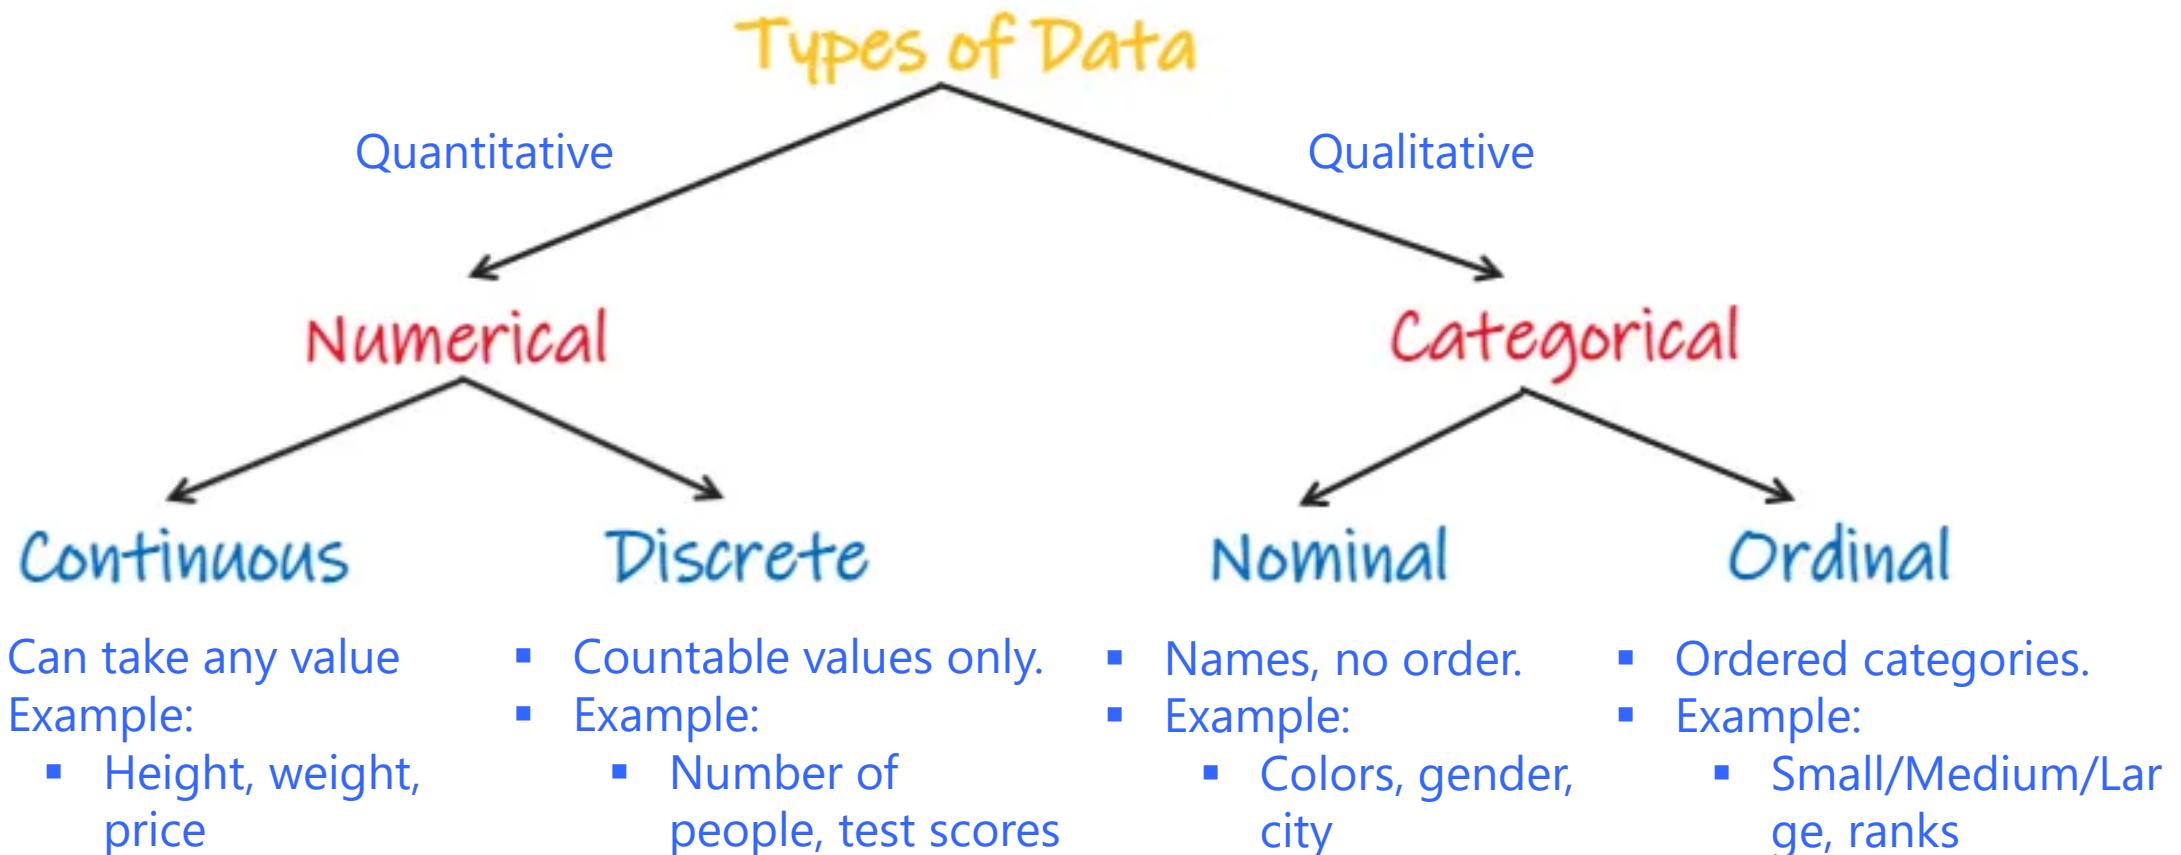
Supervised Machine Learning

- We have labeled data (input/output pairs)

Input					Output
longitude	latitude	median_age	median_income	ocean_proximity	median_house_value
-122.23	37.88	41.0	8.3252	Near Bay	452600.0
-122.22	37.86	21.0	8.3014	Near Bay	358500.0
-122.24	37.85	52.0	7.2574	Near Bay	352100.0
-122.25	37.85	52.0	5.6431	Near Bay	341300.0
-122.25	37.85	52.0	3.8462	Near Bay	342200.0

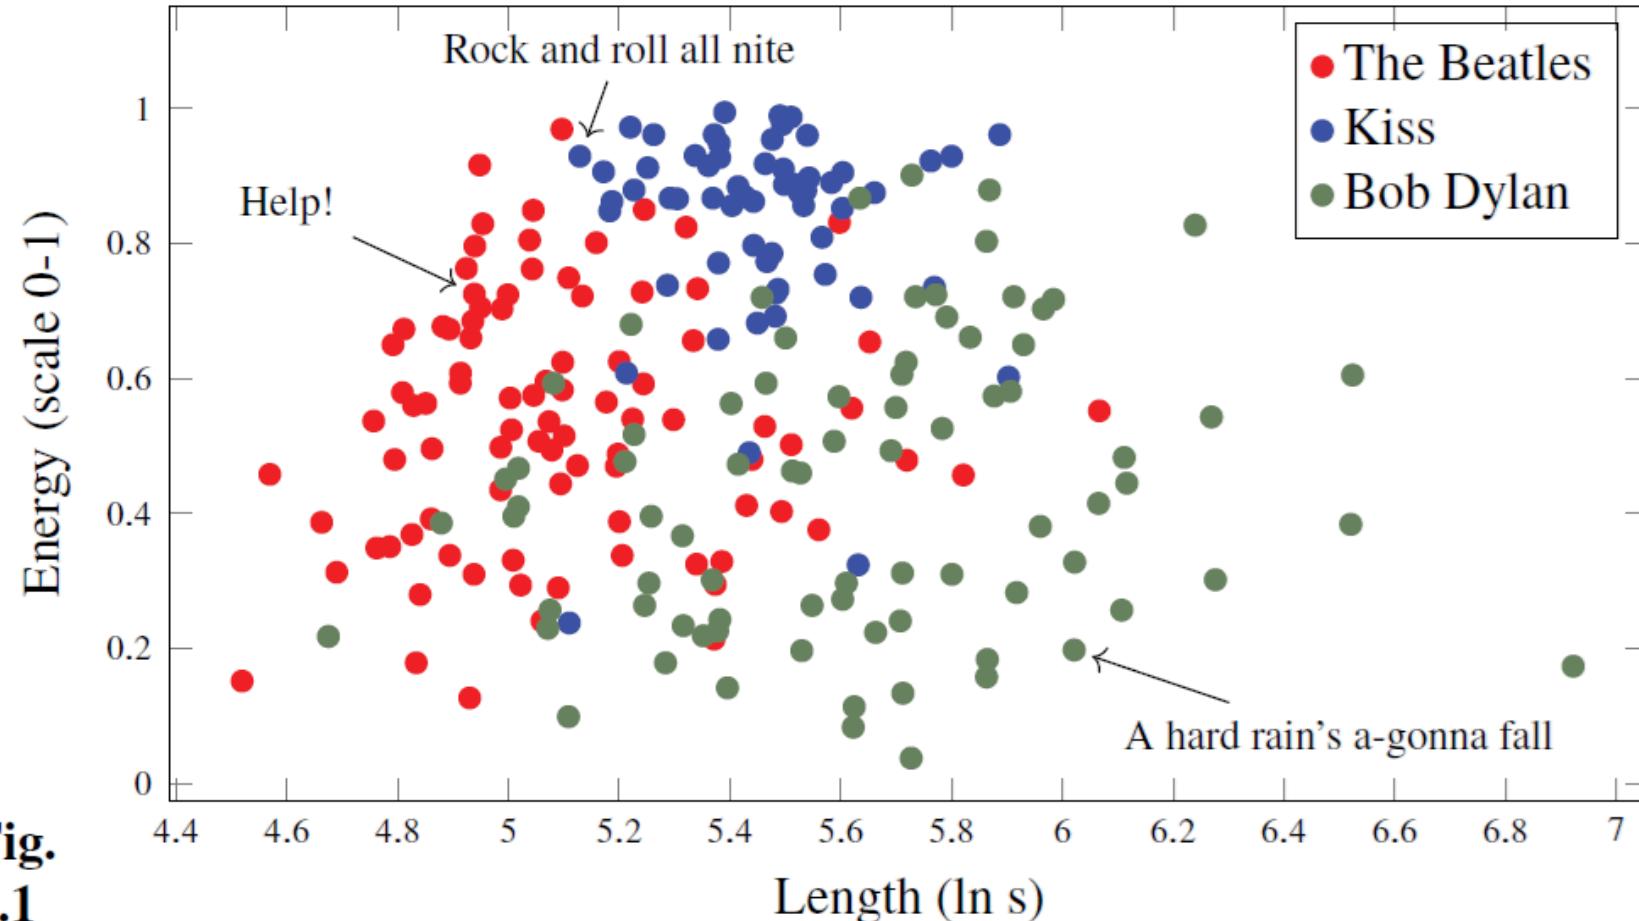
Numerical and Categorical Variables

- Two different types of variables in our data: numerical or categorical



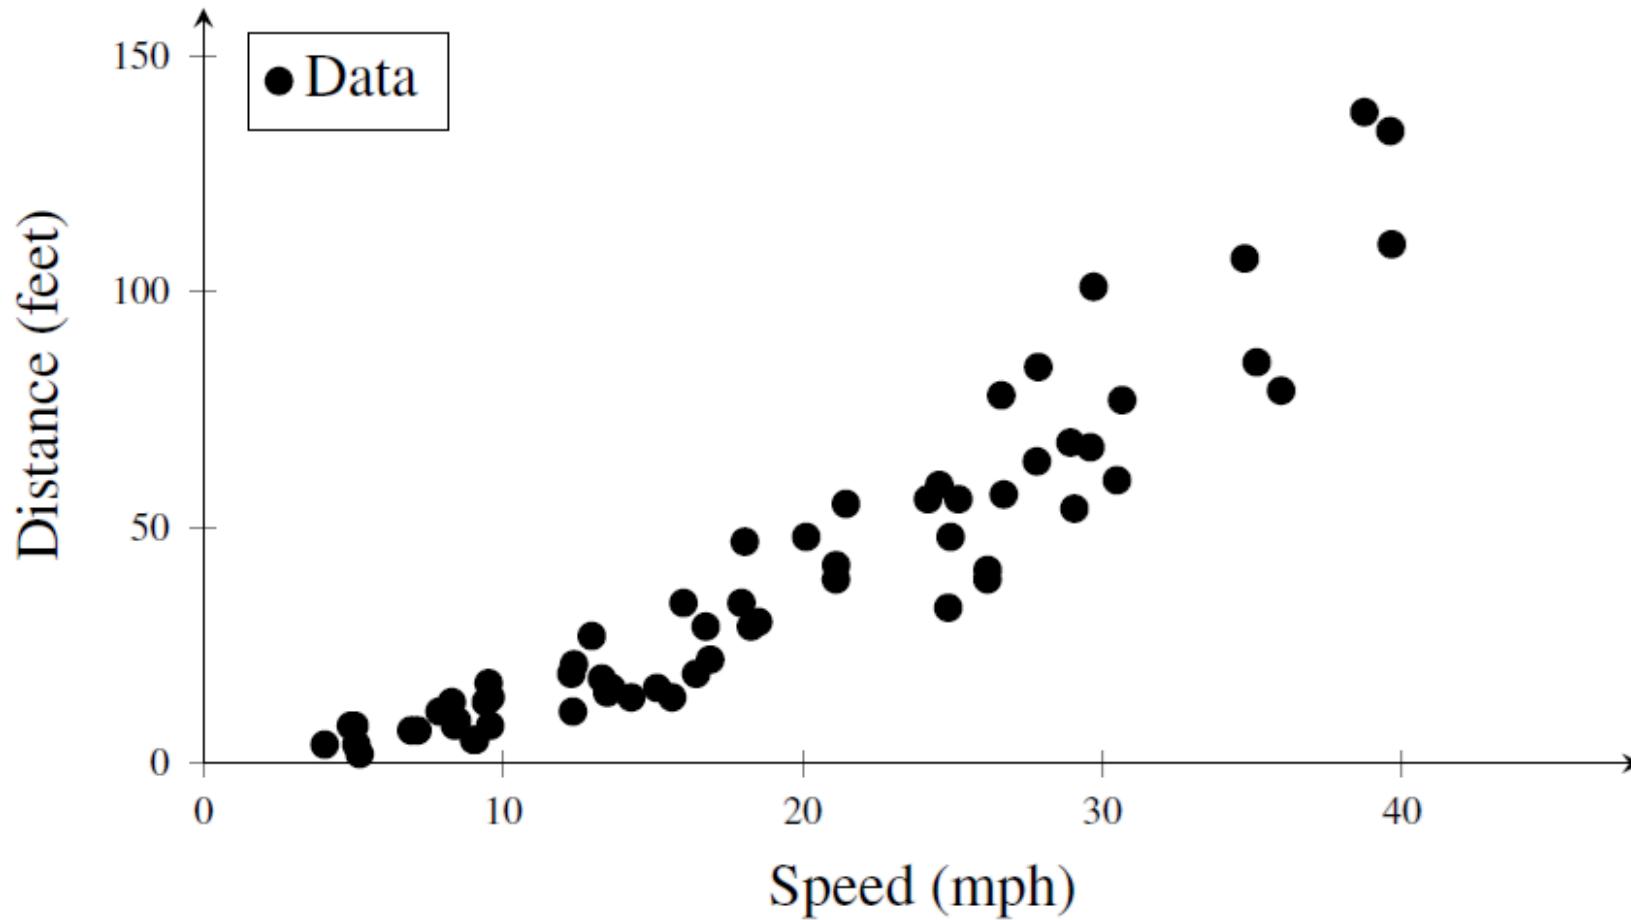
Supervised Learning - Classification

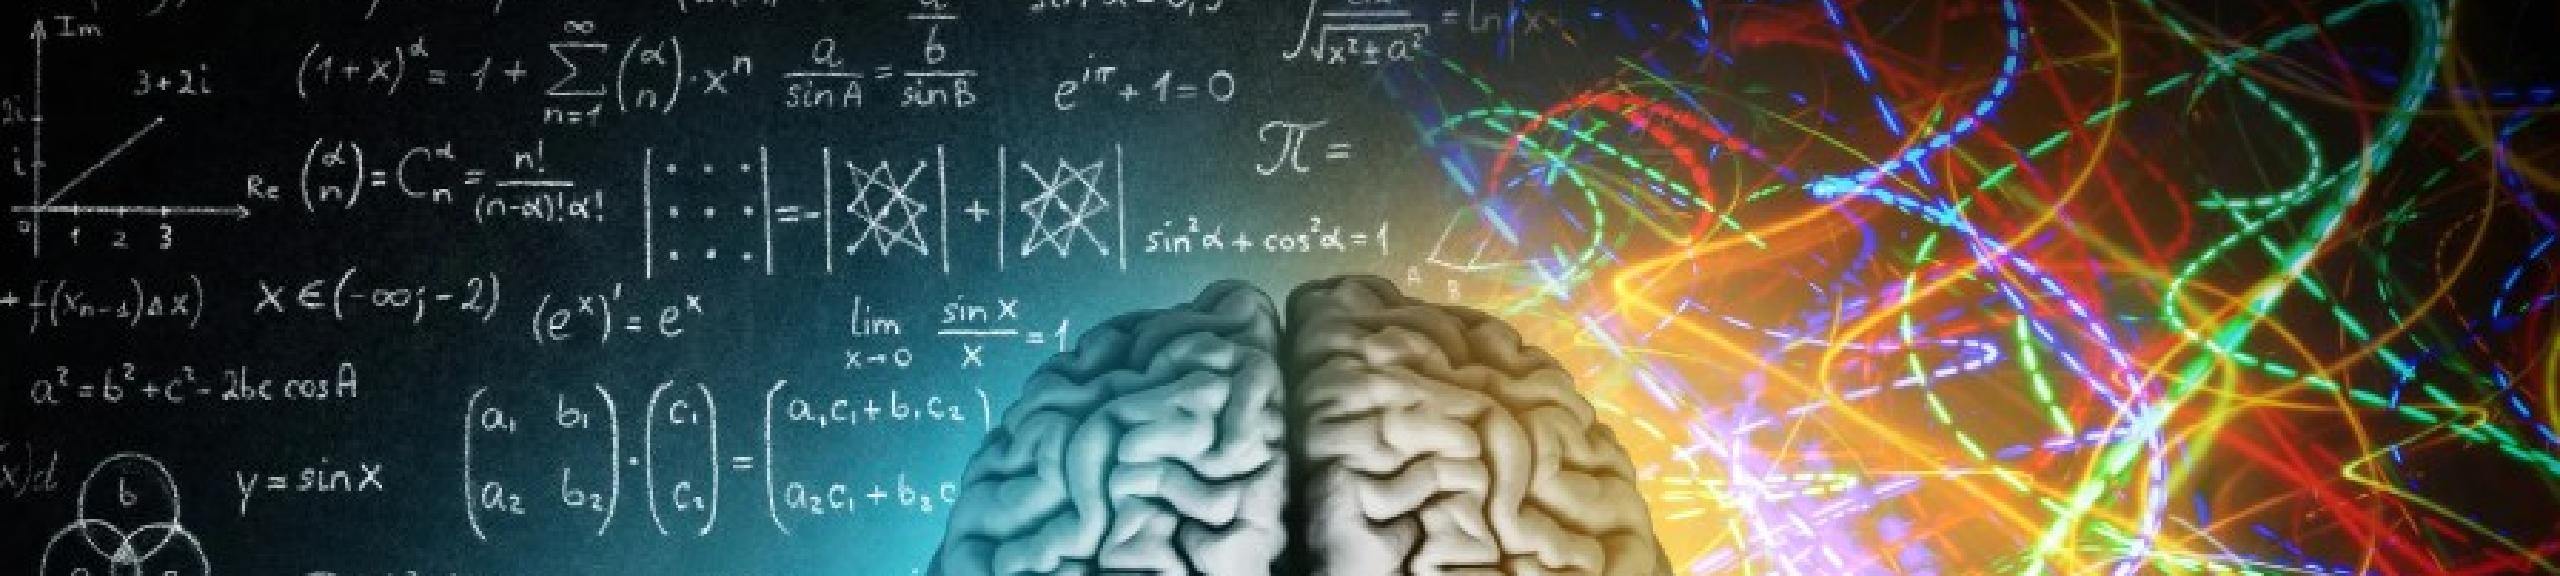
- Classifying songs
 - Length and Energy is input, and song class is output.



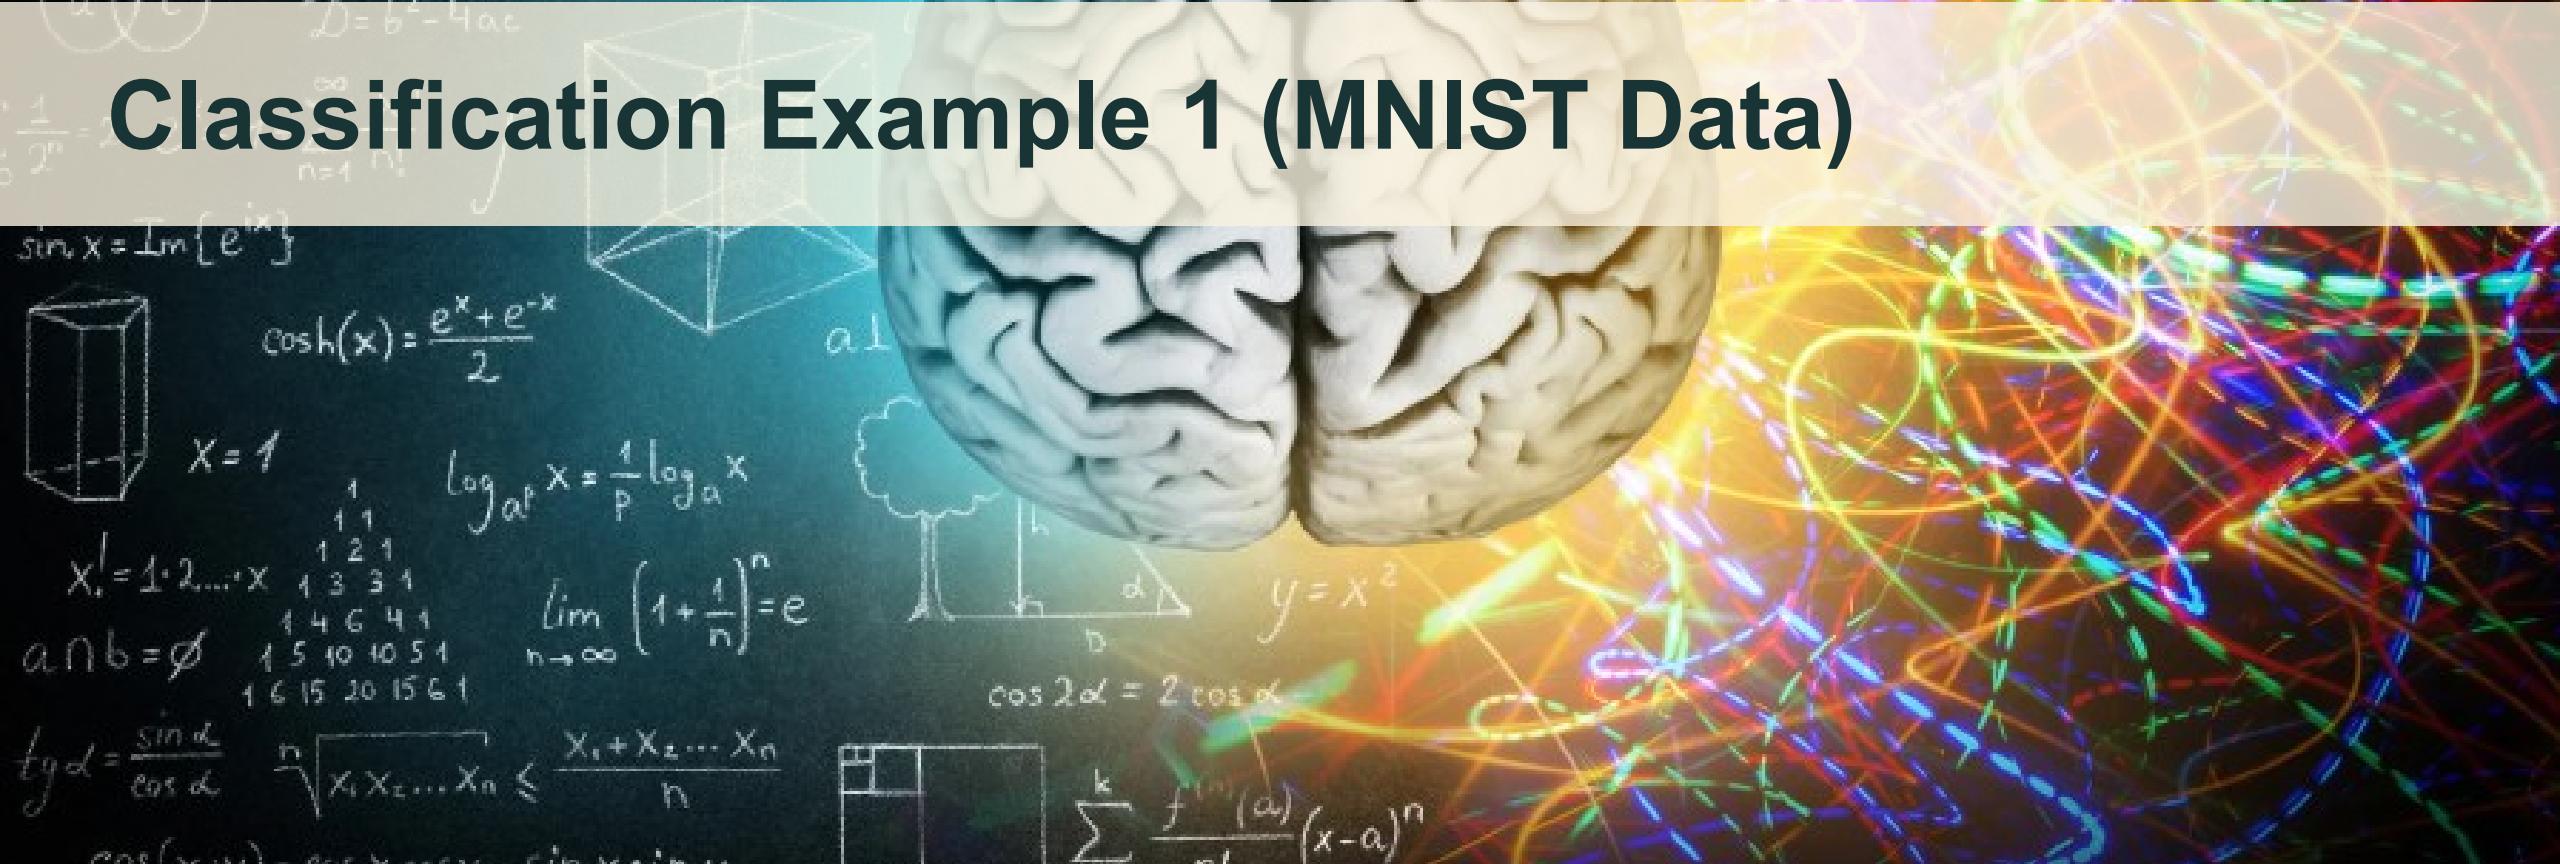
Supervised Learning - Regression

- Car stopping distance
 - Speed is input and Distance is output.





Classification Example 1 (MNIST Data)



Classification Overview

- Two main tasks in ML:
 - Regression → predict numbers (already covered)
 - Classification → predict categories (cover in this lecture)
- **MNIST** = famous dataset for digit classification (0–9)
- Called the “Hello World” of Machine Learning

MNIST Dataset

- 70,000 handwritten digit images (0–9)
- Collected from US high school students & Census Bureau staff
- Each image: 28×28 pixels = 784 features
- Pixel values: 0 = white, 255 = black



Loading MNIST in Scikit-Learn

- Use **fetch_openml()** from `sklearn.datasets`
 - `fetch_*` → download real-life datasets (e.g., `fetch_openml()`)
 - `load_*` → load small built-in toy datasets
 - `make_*` → create fake datasets for testing
- Returns data as NumPy arrays (when `as_frame=False`)
- Data (`X`) = pixel intensities
- Labels (`y`) = digit class (0–9)

```
from sklearn.datasets import fetch_openml

mnist = fetch_openml('mnist_784', as_frame=False)

X, y = mnist.data, mnist.target
```

Dataset Structure

- `X.shape` → (70000, 784) (70k images × 784 pixels)
- `y.shape` → (70000,) (label for each image)
- Pixel intensity values: 0–255

```
>>> X
array([[0., 0., 0., ..., 0., 0., 0.],
       [0., 0., 0., ..., 0., 0., 0.],
       [0., 0., 0., ..., 0., 0., 0.],
       ...,
       [0., 0., 0., ..., 0., 0., 0.],
       [0., 0., 0., ..., 0., 0., 0.],
       [0., 0., 0., ..., 0., 0., 0.]])
>>> X.shape
(70000, 784)
>>> y
array(['5', '0', '4', ..., '4', '5', '6'], dtype=object)
>>> y.shape
(70000,)
```

Viewing a Digit

- Reshape feature vector to 28×28
- Display using Matplotlib with `cmap="binary"`

```
import matplotlib.pyplot as plt

def plot_digit(image_data):
    image = image_data.reshape(28, 28)
    plt.imshow(image, cmap="binary")
    plt.axis("off")

some_digit = X[0]
plot_digit(some_digit)
plt.show()

>>> y[0]
'5'
```



looks like a 5 and label tells us the same

Viewing a Digit

- Figure shows a few more images from the MNIST dataset.

5 0 4 1 9 2 1 3 1 4
3 5 3 6 1 7 2 8 6 9
4 0 9 1 1 2 4 3 2 7
3 8 6 9 0 5 6 0 7 6
1 8 7 9 3 9 8 5 9 3
3 0 7 4 9 8 0 9 4 1
4 4 6 0 4 5 6 1 0 0
1 7 1 6 3 0 2 1 1 7
8 0 2 6 7 8 3 9 0 4
6 7 4 6 8 0 7 8 3 1

Train Test Split

- MNIST dataset returned by `fetch_openml()` is **already split** into a **training set** (the first 60,000 images) and a **test set** (the last 10,000 images)

```
X_train, X_test, y_train, y_test = X[:60000], X[60000:], y[:60000], y[60000:]
```

Training Data

Test Data

image



label

'5'

'0'

'4'

'1'

• • •

'9'

'2'

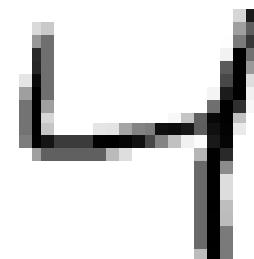
Training a Binary Classifier

- A classifier that **predicts two classes** only.
- Our task: detect **digit 5** (class 1) or **not 5** (class 0)
- Example: “5-detector” model

✓ (is 5)



✗ (not 5)



```
X_train, X_test, y_train, y_test = X[:60000], X[60000:], y[:60000], y[60000:]
```

Choose a Classifier (**SGDClassifier**)

- **SGDClassifier** = Stochastic Gradient Descent Classifier

- Handles very large datasets well
- Learns one example at a time → good for online learning

```
from sklearn.linear_model import SGDClassifier  
  
sgd_clf = SGDClassifier(random_state=42)  
sgd_clf.fit(X_train, y_train_5)  
  
some_digit = X[0]  
  
>>> sgd_clf.predict([some_digit])  
array([ True])
```

Performance Measures

- **Classifier** evaluation is **harder** than regression
- Many metrics to learn (accuracy, confusion matrix, etc.)
- First: **measure accuracy** with cross-validation
 - Cross-validation = split training set into k folds
 - Train/test k times → more reliable score

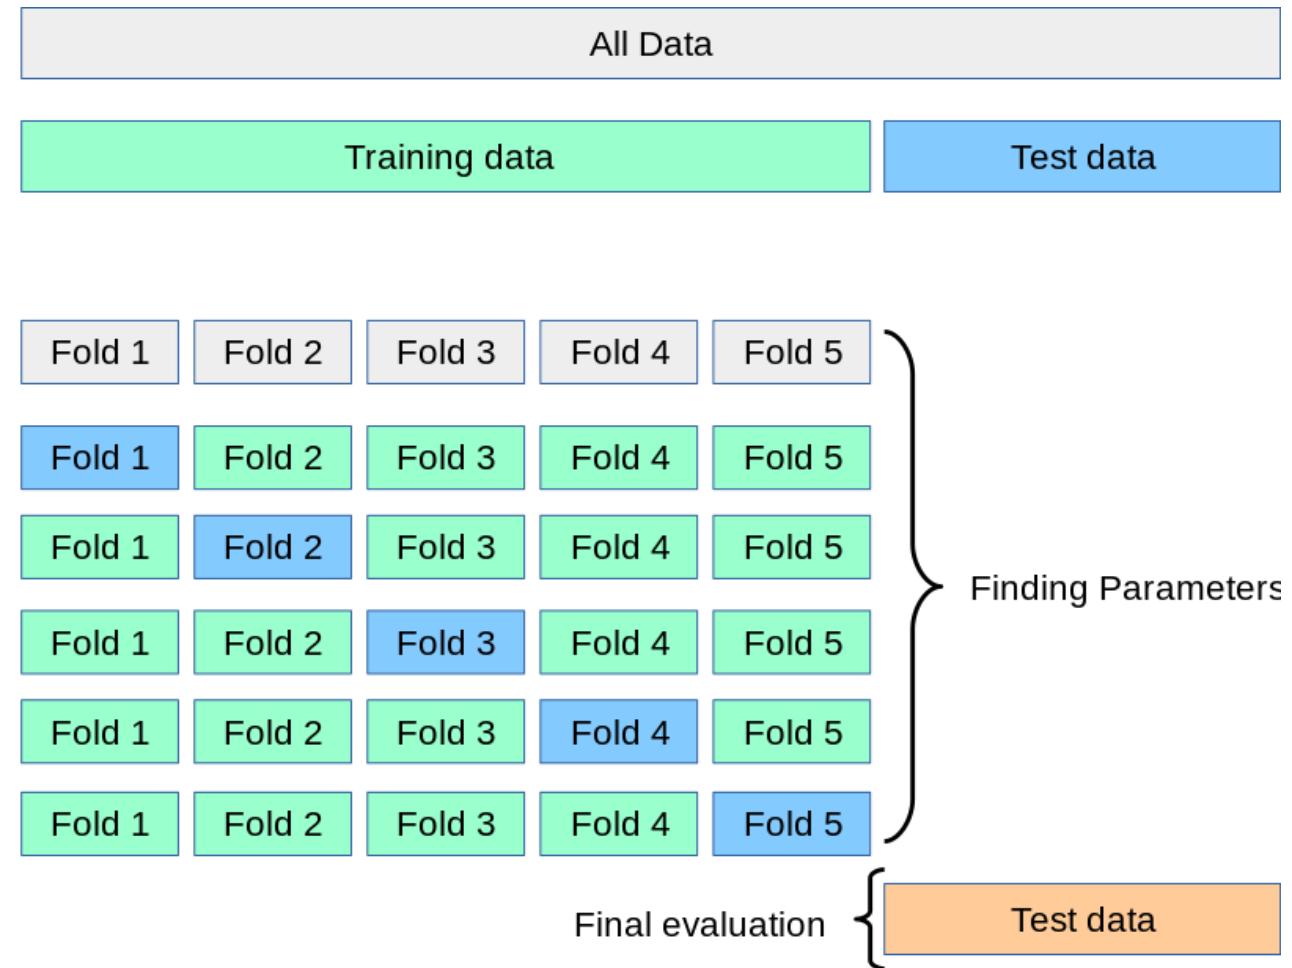
```
>>> from sklearn.model_selection import cross_val_score  
>>> cross_val_score(sgd_clf, X_train, y_train_5, cv=3, scoring="accuracy")  
array([0.95035, 0.96035, 0.9604 ])
```

Wow! Above 95% accuracy

Performance Measures

Cross Validation Concept

- In k-fold cross-validation, the data is divided into k subsets, and the model is trained and validated k times.
- Each time, one of the k subsets is used as the test set, and the remaining k-1 subsets form the training set.
- The final performance is the average of the k trials.



The Dummy Classifier Trick

- Always predicts most common class
- Here: always “not 5” → 90% accuracy (because only about 10% of the images are 5s)
- Shows why accuracy can mislead in skewed data.

```
from sklearn.dummy import DummyClassifier

dummy_clf = DummyClassifier()
dummy_clf.fit(X_train, y_train_5)
print(any(dummy_clf.predict(X_train))) # prints False: no 5s detected

>>> cross_val_score(dummy_clf, X_train, y_train_5, cv=3, scoring="accuracy")
array([0.90965, 0.90965, 0.90965])
```

Manual Cross-Validation (StratifiedKFold)

- `StratifiedKFold` → keeps class proportions same in each fold
- More control over CV process
- Fresh model each fold (`clone()`)

```
from sklearn.model_selection import StratifiedKFold
from sklearn.base import clone

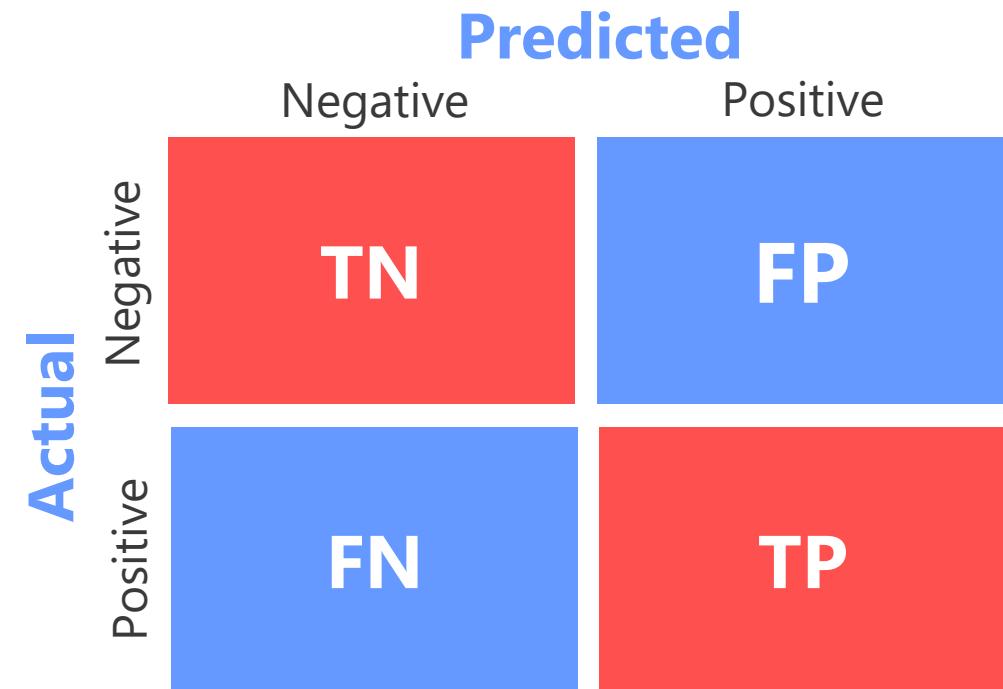
skfolds = StratifiedKFold(n_splits=3) # add shuffle=True if the dataset is
# not already shuffled
for train_index, test_index in skfolds.split(X_train, y_train_5):
    clone_clf = clone(sgd_clf)
    X_train_folds = X_train[train_index]
    y_train_folds = y_train_5[train_index]
    X_test_fold = X_train[test_index]
    y_test_fold = y_train_5[test_index]

    clone_clf.fit(X_train_folds, y_train_folds)
    y_pred = clone_clf.predict(X_test_fold)
    n_correct = sum(y_pred == y_test_fold)
    print(n_correct / len(y_pred)) # prints 0.95035, 0.96035, and 0.9604
```

Confusion Matrix

- Confusion matrix summarizes classifier's accuracy.
- Matrix shows correct and incorrect classifications.
- Rows = actual class, columns = predicted class
- Obtain using `confusion_matrix()` in scikit-learn.

		Predicted	
		No diabetes	Diabetes
Actual	No diabetes	201	85
	Diabetes	25	2689



Making the Confusion Matrix

- We need predictions before making confusion matrix
- Use `cross_val_predict()` → makes clean predictions (model never sees test fold before predicting)
- Use `confusion_matrix()` with true labels and predicted labels

```
from sklearn.model_selection import cross_val_predict

y_train_pred = cross_val_predict(sgd_clf, X_train, y_train_5, cv=3)

>>> from sklearn.metrics import confusion_matrix
>>> cm = confusion_matrix(y_train_5, y_train_pred)
>>> cm
array([[53892,    687],
       [ 1891,  3530]])
```

Understanding the Numbers

- True Negatives (TN) → correct “Not 5” predictions (53,892)
- False Positives (FP) → “Not 5” predicted as “5” (687) → Type I error
- False Negatives (FN) → “5” predicted as “Not 5” (1,891) → Type II error
- True Positives (TP) → correct “5” predictions (3,530)

		Predicted
		Negative
Actual	Negative	TN
	Positive	FP
		TP
		FN

		Predicted
		Not “5”
Actual	Not “5”	53892
	“5”	687
		“5”
		1891
		3530

Perfect Classifier Example

- Only TP & TN (no FP, no FN)
- Confusion matrix has numbers only on main diagonal

```
>>> y_train_perfect_predictions = y_train_5 # pretend we reached perfection  
>>> confusion_matrix(y_train_5, y_train_perfect_predictions)  
array([[54579,      0],  
       [      0,  5421]])
```

		Predicted
		Negative
Actual	Negative	TN
	Positive	FP
Positive	Negative	FN
Positive	Positive	TP

		Predicted
		Not "5"
Actual	Not "5"	54579
	"5"	0
"5"	Not "5"	0
"5"	"5"	5421

Accuracy – Proportion of Correct Prediction

- Proportion of **all correct predictions** in the dataset.
- The **overall accuracy** of a **classifier** is estimated by

$$Accuracy = \frac{TP + TN}{TP + TN + FP + FN}$$

$$Accuracy = \frac{3530 + 53892}{3530 + 53892 + 687 + 1891} = 0.914$$

		Predicted	
		Negative	Positive
Actual	Negative	TN	FP
	Positive	FN	TP

		Predicted	
		Not "5"	"5"
Actual	Not "5"	53892	687
	"5"	1891	3530

Precision – Positive Prediction Accuracy

- Measures how many **predicted positives** are **correct**.
- Assesses the exactness or quality of the classifier
- Out of 4217 predicted as positive ("5"), 3530 were correct.
- Not useful alone → ignores missed positives

$$Precision = \frac{TP}{TP + FP}$$

$$Precision = \frac{3530}{3530 + 687} = 0.837$$

		Predicted	
		Negative	Positive
Actual	Negative	TN	FP
	Positive	FN	TP
		Predicted	
		Not "5"	"5"
Actual	Not "5"	53892	687
	"5"	1891	3530

Recall – True Positive Rate

- Measures how many **real positives** were found
- Ratio of true positives to total (actual) positives in the data.
- Also called Sensitivity or True Positive Rate
- Of 5421 actual "5", 3530 were predicted as "5".

$$Recall = \frac{TP}{TP + FN}$$

$$Recall = \frac{3530}{3530 + 1891} = 0.651$$

		Predicted	
		Negative	Positive
Actual	Negative	TN	FP
	Positive	FN	TP
		Predicted	
		Not "5"	"5"
Actual	Not "5"	53892	687
	"5"	1891	3530

F1-Score

- Harmonic mean of precision and recall
- Combine precision and recall into a single metric
- Use when we need a single metric to compare two classifiers

$$F1\text{-}Score = \frac{2 \times (Precision \times Recall)}{(Precision + Recall)}$$

$$F1\text{-}Score = \frac{2 \times (0.837 \times 0.651)}{(0.837 + 0.651)} = 0.732$$

		Predicted	
		Negative	Positive
Actual	Negative	TN	FP
	Positive	FN	TP
		Predicted	
		Not "5"	"5"
Actual	Not "5"	53892	687
	"5"	1891	3530

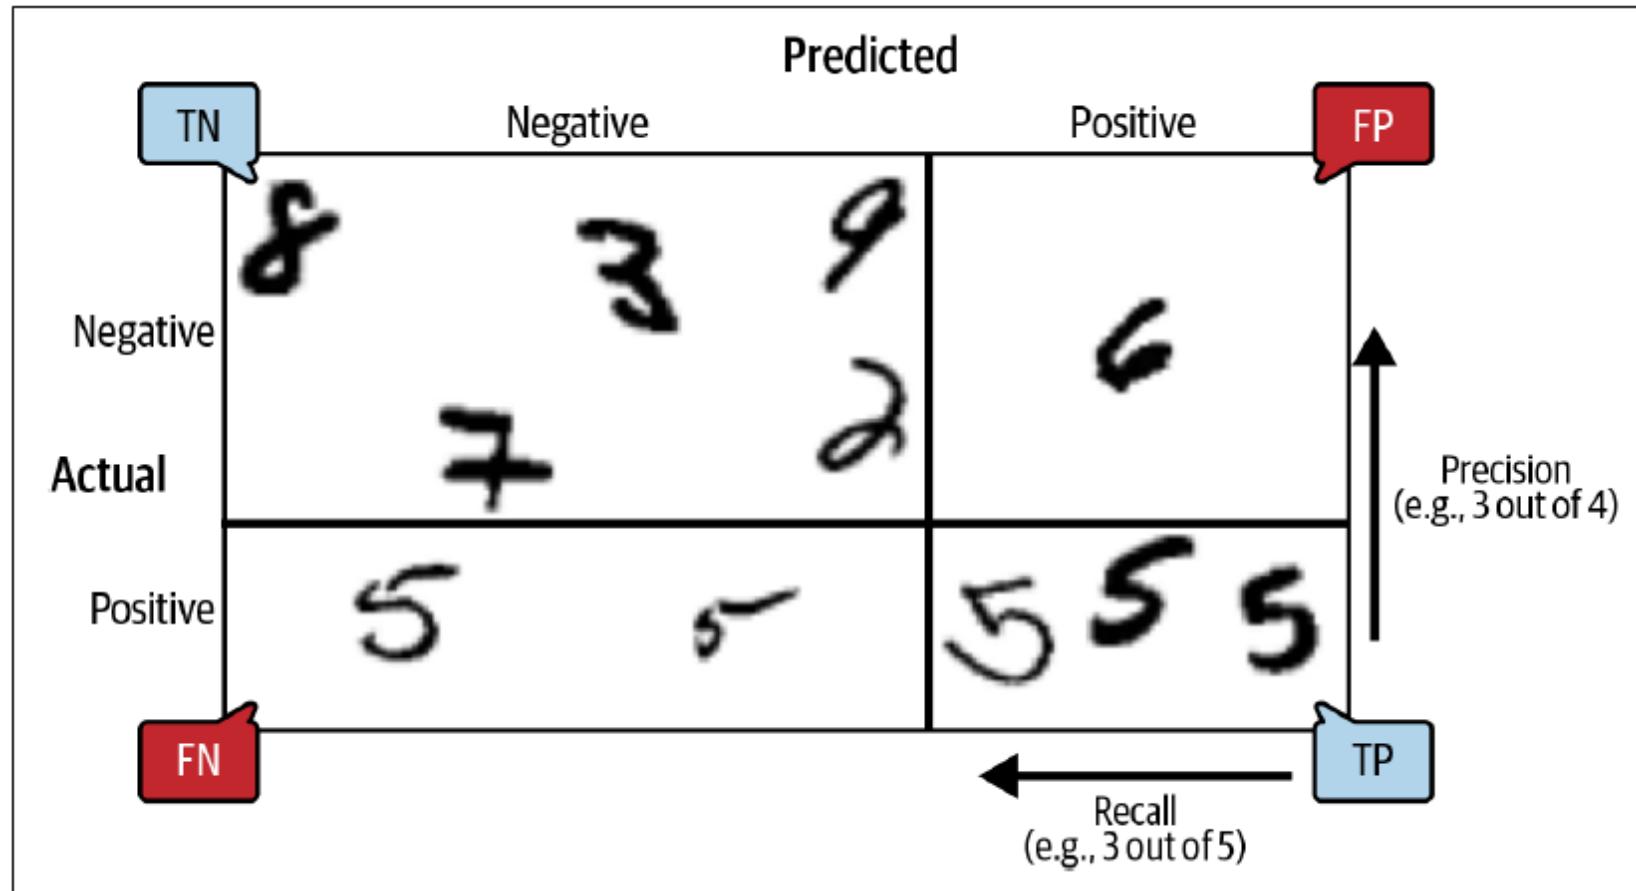
Precision, Recall and F1-Score in Sklearn

- Scikit-Learn provides several functions to compute classifier metrics, including precision, recall and F1-Score.

```
>>> from sklearn.metrics import precision_score, recall_score  
>>> precision_score(y_train_5, y_train_pred) # == 3530 / (687 + 3530)  
0.8370879772350012  
>>> recall_score(y_train_5, y_train_pred) # == 3530 / (1891 + 3530)  
0.6511713705958311  
>>> from sklearn.metrics import f1_score  
>>> f1_score(y_train_5, y_train_pred)  
0.7325171197343846
```

Confusion Matrix - Summary

- If you are confused about the confusion matrix, Figure below may help.

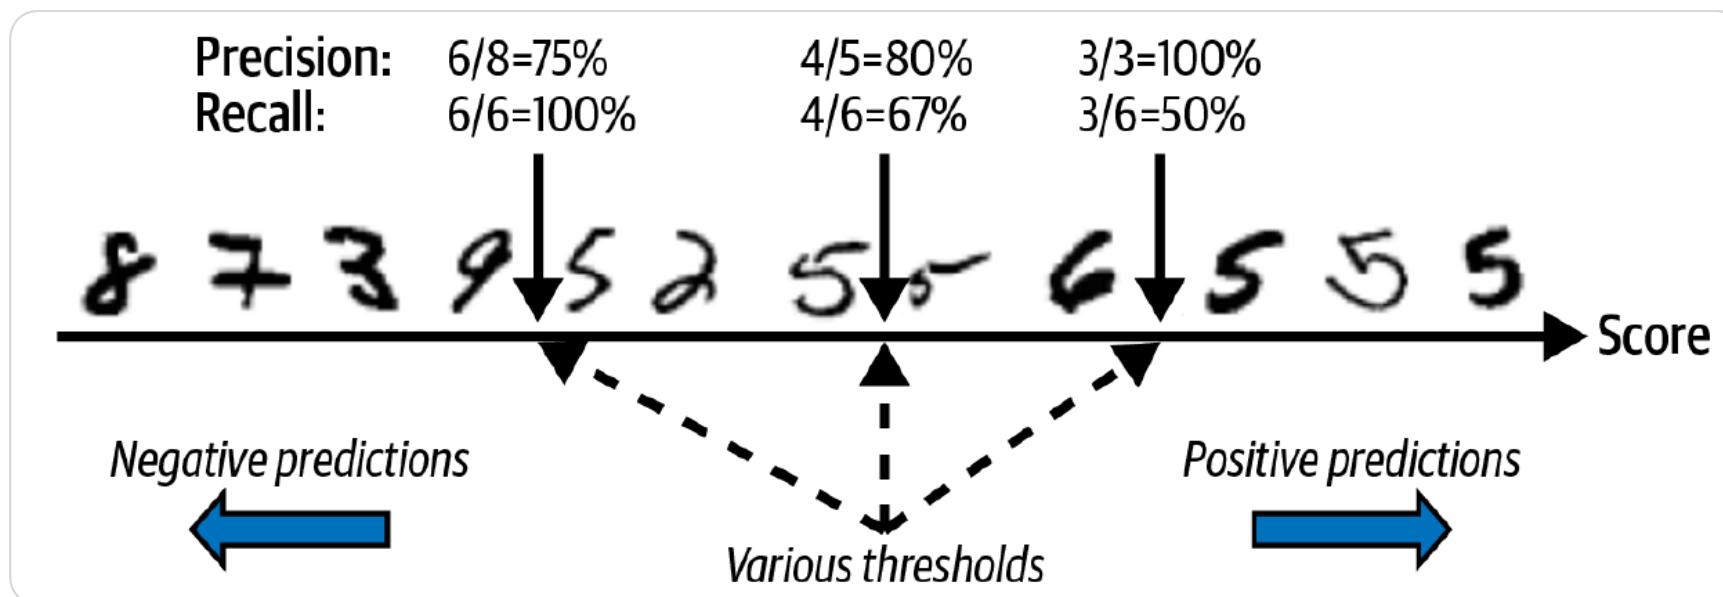


Precision/Recall Trade-off

- Some problems need high precision (how many predicted positives are correct)
 - Example: Safe videos for kids
 - Important to avoid bad content, even if we miss good ones
- Some problems need high recall
 - Example: Detecting shoplifters
 - Okay to raise false alarms, but don't miss real ones
- You can't increase both at the same time. This is called a trade-off
 -  High Precision →  Low Recall
 -  High Recall →  Low Precision

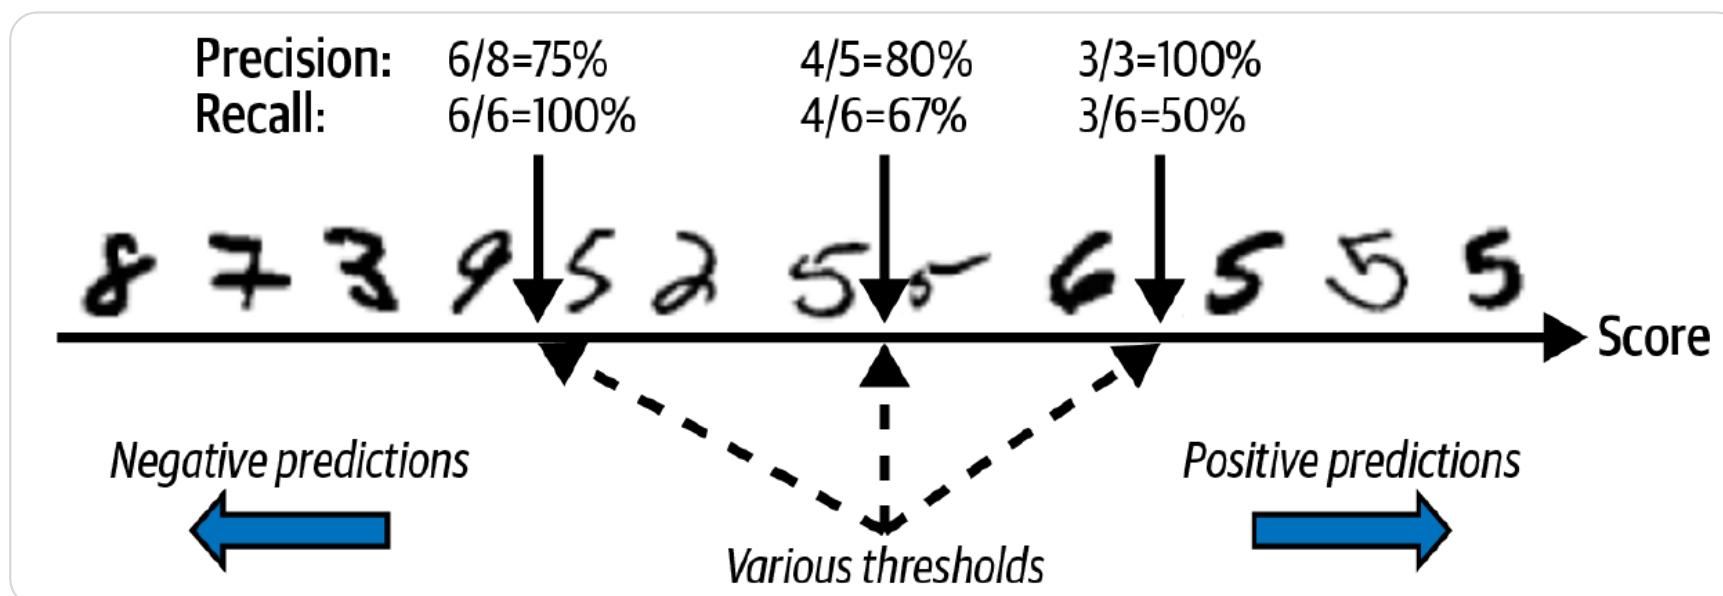
Precision/Recall Trade-off – Example

- Classifier gives each instance a score
- If score > threshold → predict positive (5)
- If score < threshold → predict negative (not-5)
- Moving threshold changes precision and recall.



Precision/Recall Trade-off – Example

- Threshold at center:
 - 4 true positives (actual 5s)
 - 1 false positive (6 wrongly predicted as 5)
 - Precision = 80% ($4/5$), Recall = 67% ($4/6$)
- Raise threshold → precision = 100%, recall = 50%
- Lower threshold → recall increases, precision decreases



Controlling Threshold in Code

- By default, SGDClassifier uses threshold = 0
- You can access decision score using:

```
>>> y_scores = sgd_clf.decision_function([some_digit])
>>> y_scores
array([2164.22030239])
```

- Then apply any threshold:

```
>>> threshold = 0
>>> y_some_digit_pred = (y_scores > threshold)
array([ True])
>>> threshold = 3000
>>> y_some_digit_pred = (y_scores > threshold)
>>> y_some_digit_pred
array([False])
```

How do You Decide which Threshold to Use?

- First, get decision scores for all samples:

```
y_scores = cross_val_predict(sgd_clf, X_train, y_train_5, cv=3,  
                             method="decision_function")
```

- Then compute precision, recall, and thresholds:

```
from sklearn.metrics import precision_recall_curve
```

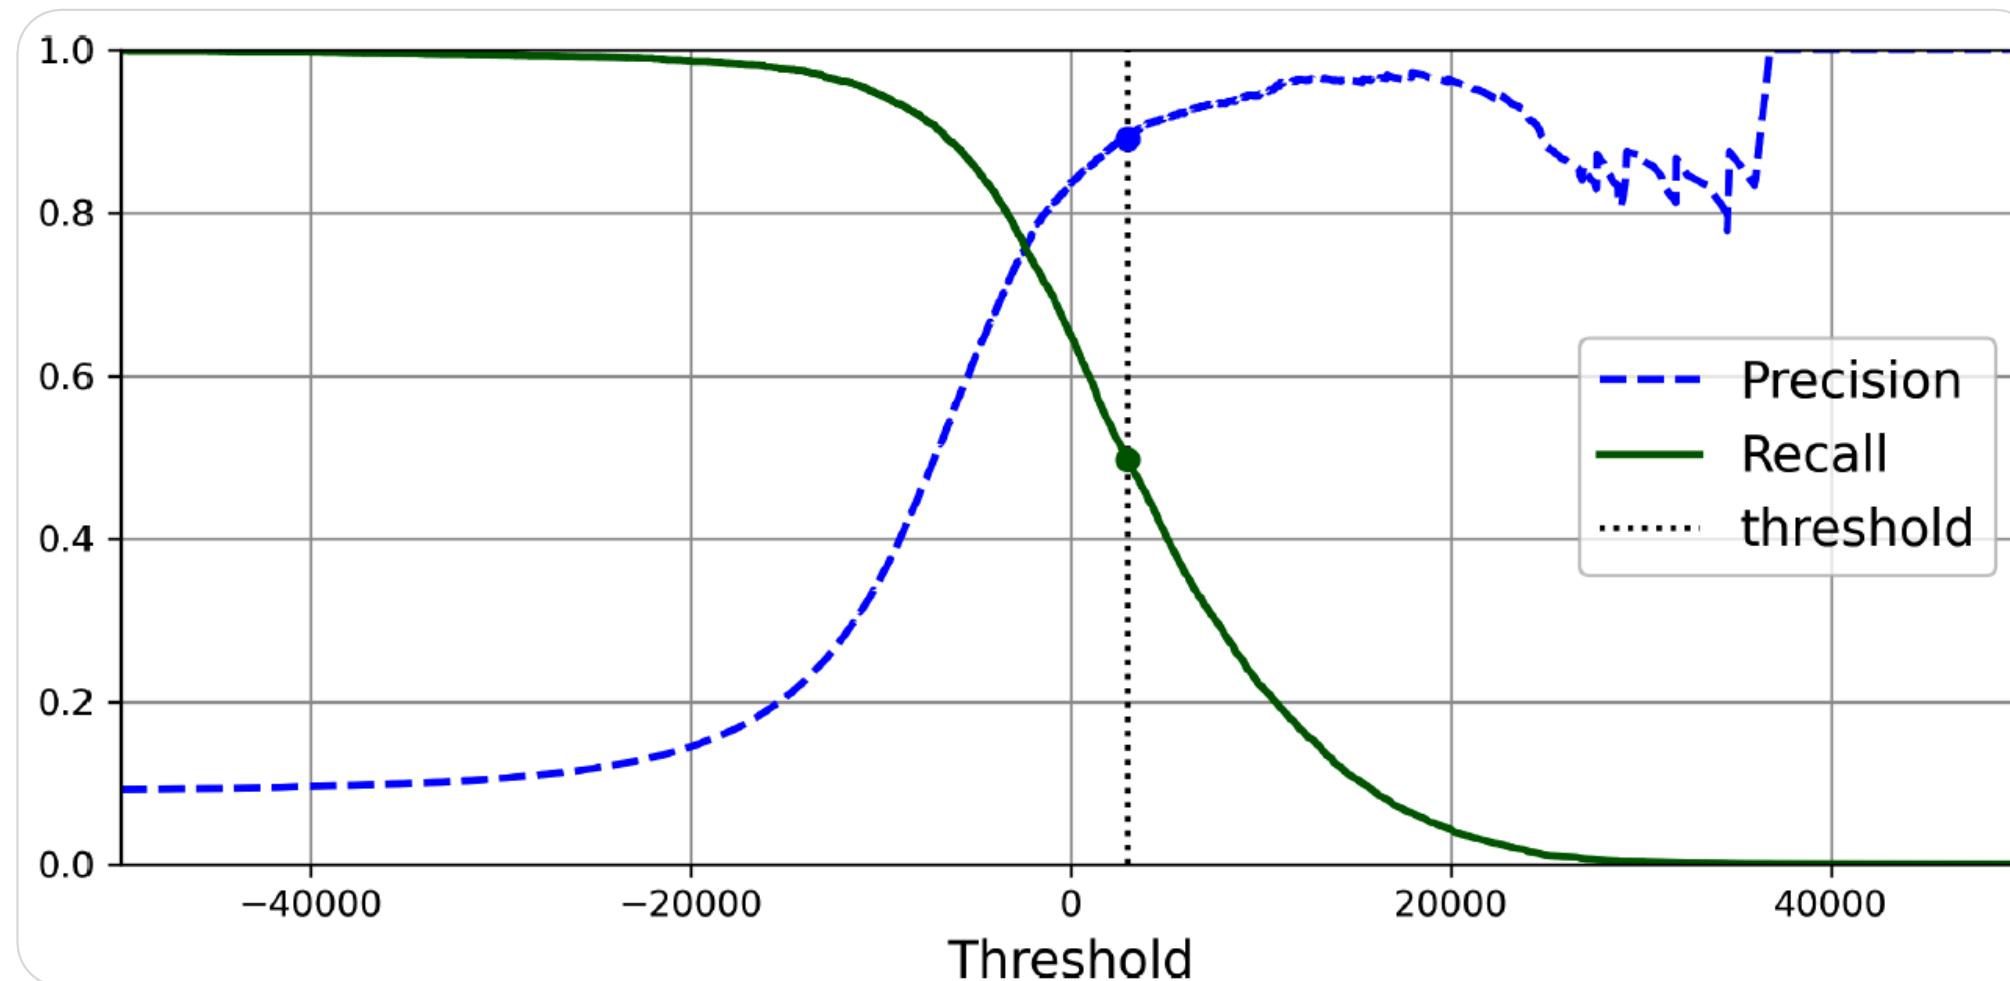
```
precisions, recalls, thresholds = precision_recall_curve(y_train_5, y_scores)
```

- Plot curves:

```
plt.plot(thresholds, precisions[:-1], "b--", label="Precision", linewidth=2)  
plt.plot(thresholds, recalls[:-1], "g-", label="Recall", linewidth=2)  
plt.vlines(threshold, 0, 1.0, "k", "dotted", label="threshold")  
[...] # beautify the figure: add grid, legend, axis, labels, and circles  
plt.show()
```

How do You Decide which Threshold to Use?

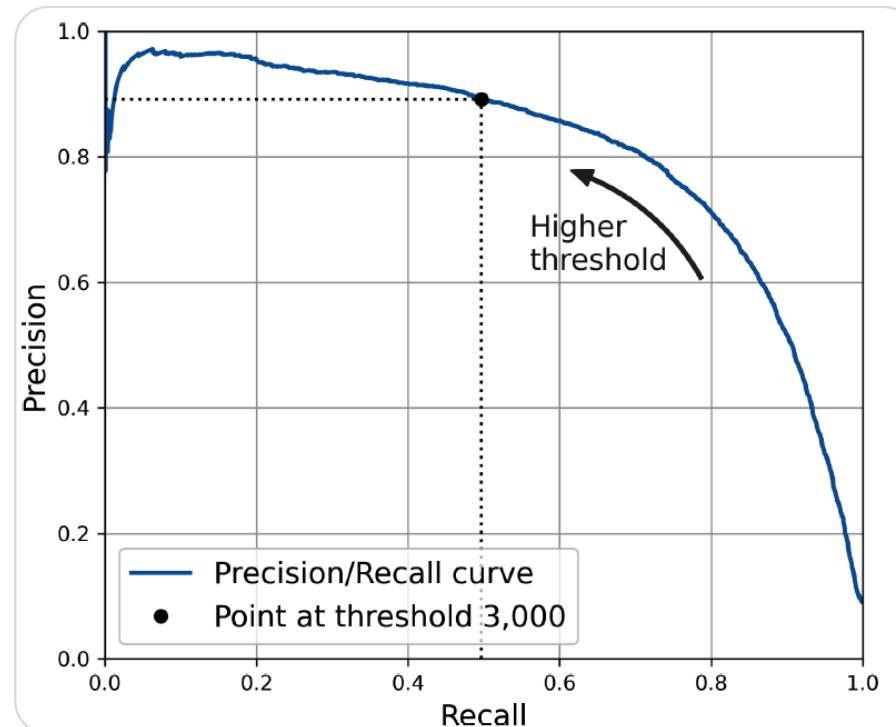
- At threshold $\approx 3000 \rightarrow$ precision $\approx 90\%$, recall $\approx 50\%$



Precision vs Recall Curve

- Plot precision directly against recall
- Shows trade-off clearly
- Precision drops fast after ~80% recall
- Good idea: pick trade-off before sharp drop (e.g., 60%)

```
plt.plot(recalls, precisions, linewidth=2, label="Precision/Recall curve")
[...] # beautify the figure: add labels, grid, legend, arrow, and text
plt.show()
```



Choosing a Target Precision

- Can search for threshold giving at least 90% precision
- Use NumPy argmax() to find correct threshold

```
>>> idx_for_90_precision = (precisions >= 0.90).argmax()  
>>> threshold_for_90_precision = thresholds[idx_for_90_precision]  
>>> threshold_for_90_precision  
3370.0194991439557
```

- Then make predictions:

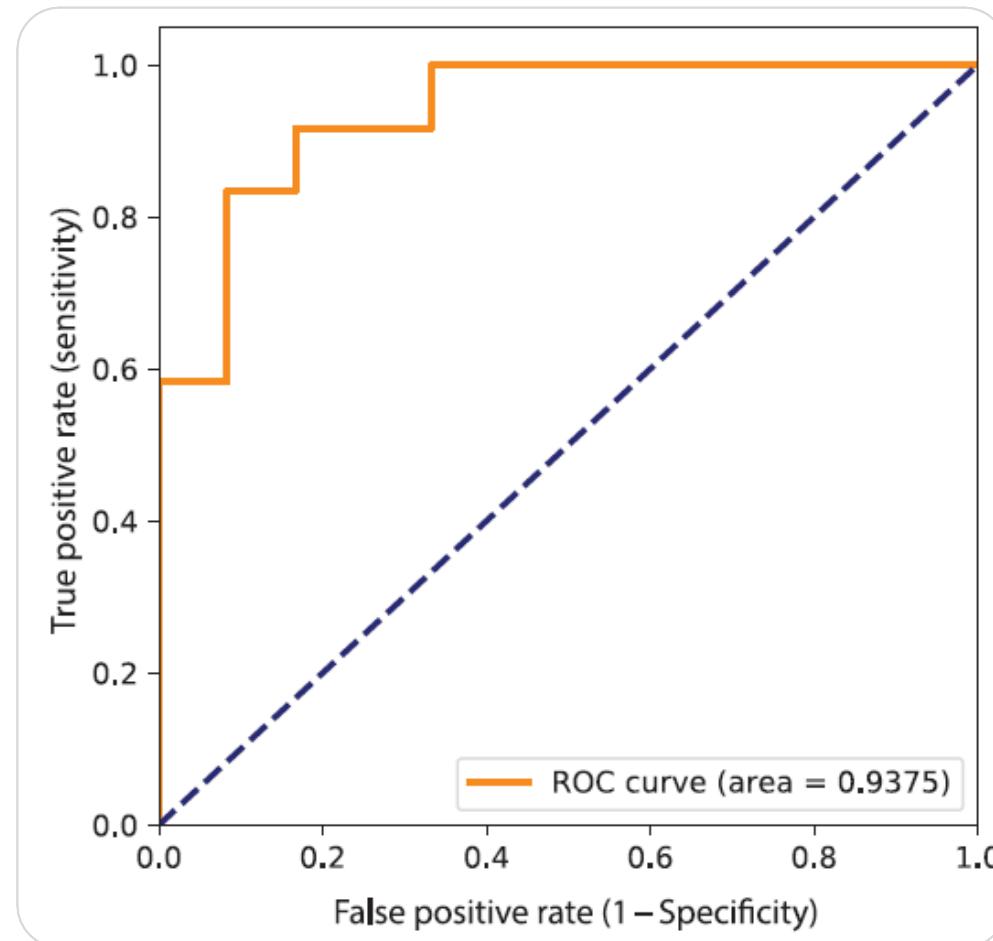
```
y_train_pred_90 = (y_scores >= threshold_for_90_precision)
```

- Result: precision \approx 90%, recall \approx 48%

```
>>> precision_score(y_train_5, y_train_pred_90)  
0.9000345901072293  
>>> recall_at_90_precision = recall_score(y_train_5, y_train_pred_90)  
>>> recall_at_90_precision  
0.4799852425751706
```

ROC = Receiver Operating Characteristic Curve

- Plots True Positive Rate (Recall) vs False Positive Rate
- FPR = negatives wrongly classified as positives



Plotting the ROC Curve

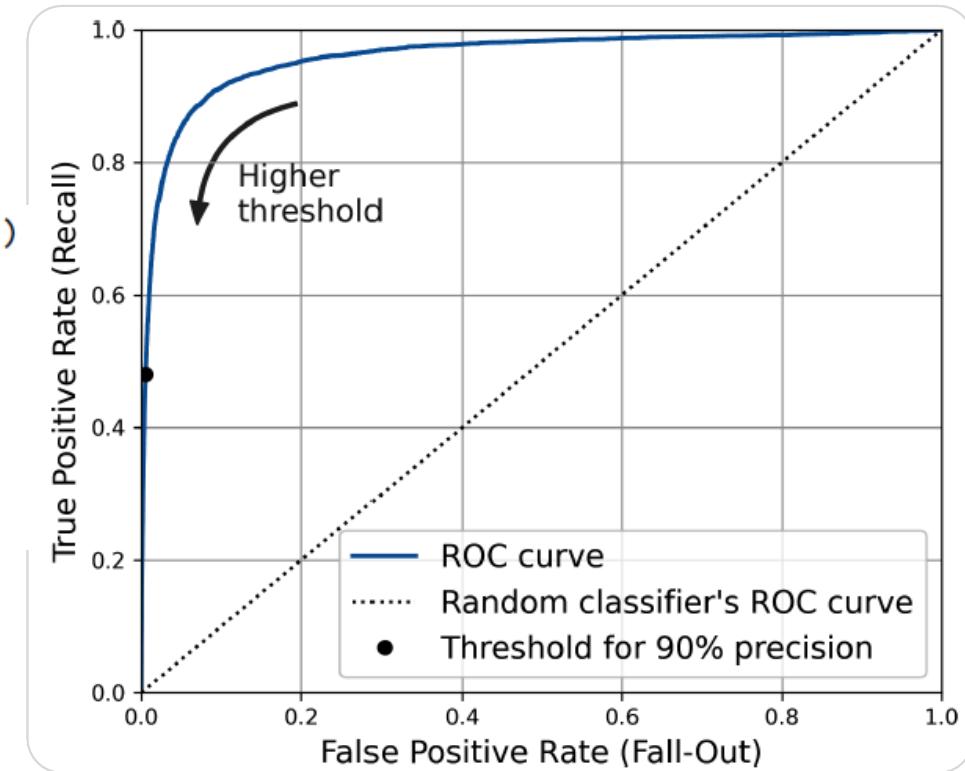
- Use `roc_curve()` to get FPR, TPR, thresholds
- Then plot FPR (x-axis) vs TPR (y-axis)
- Dotted diagonal = random classifier baseline
- Good models stay near top-left corner

```
from sklearn.metrics import roc_curve

fpr, tpr, thresholds = roc_curve(y_train_5, y_scores)

idx_for_threshold_at_90 = (thresholds <= threshold_for_90_precision).argmax()
tpr_90, fpr_90 = tpr[idx_for_threshold_at_90], fpr[idx_for_threshold_at_90]

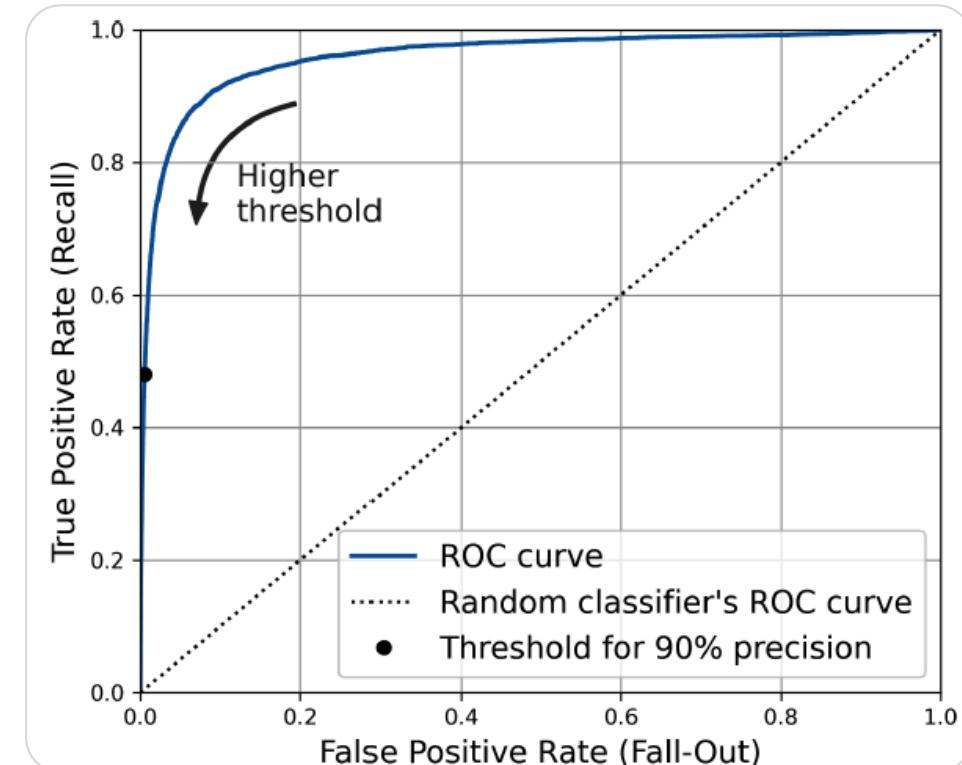
plt.plot(fpr, tpr, linewidth=2, label="ROC curve")
plt.plot([0, 1], [0, 1], 'k:', label="Random classifier's ROC curve")
plt.plot([fpr_90], [tpr_90], "ko", label="Threshold for 90% precision")
[...] # beautify the figure: add labels, grid, legend, arrow, and text
plt.show()
```

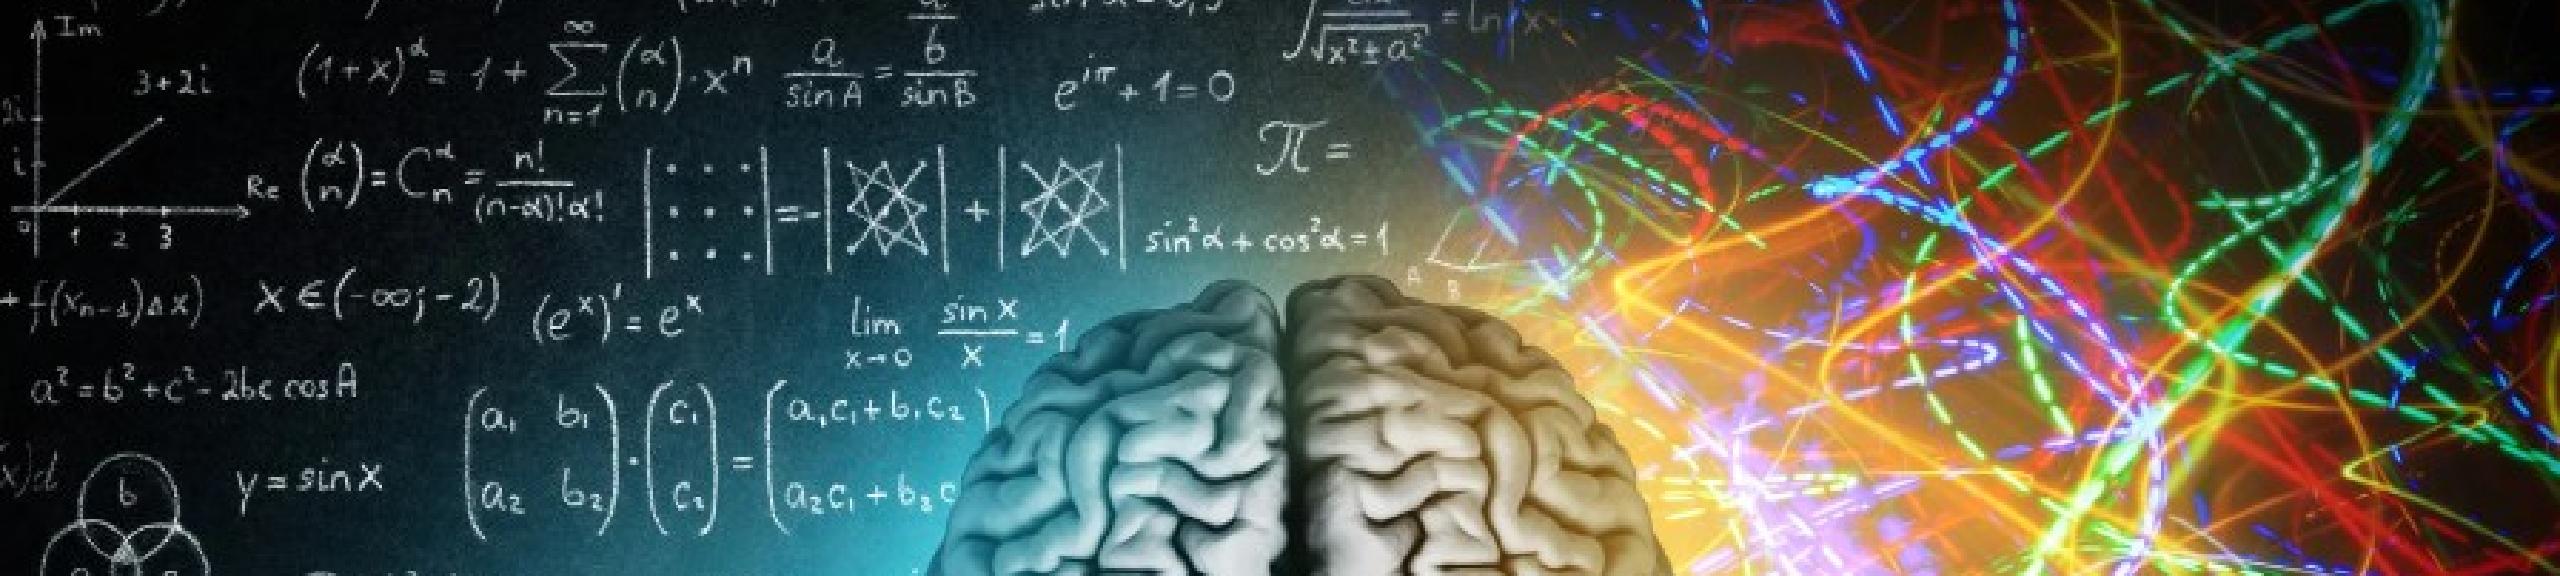


Area Under the Curve (AUC)

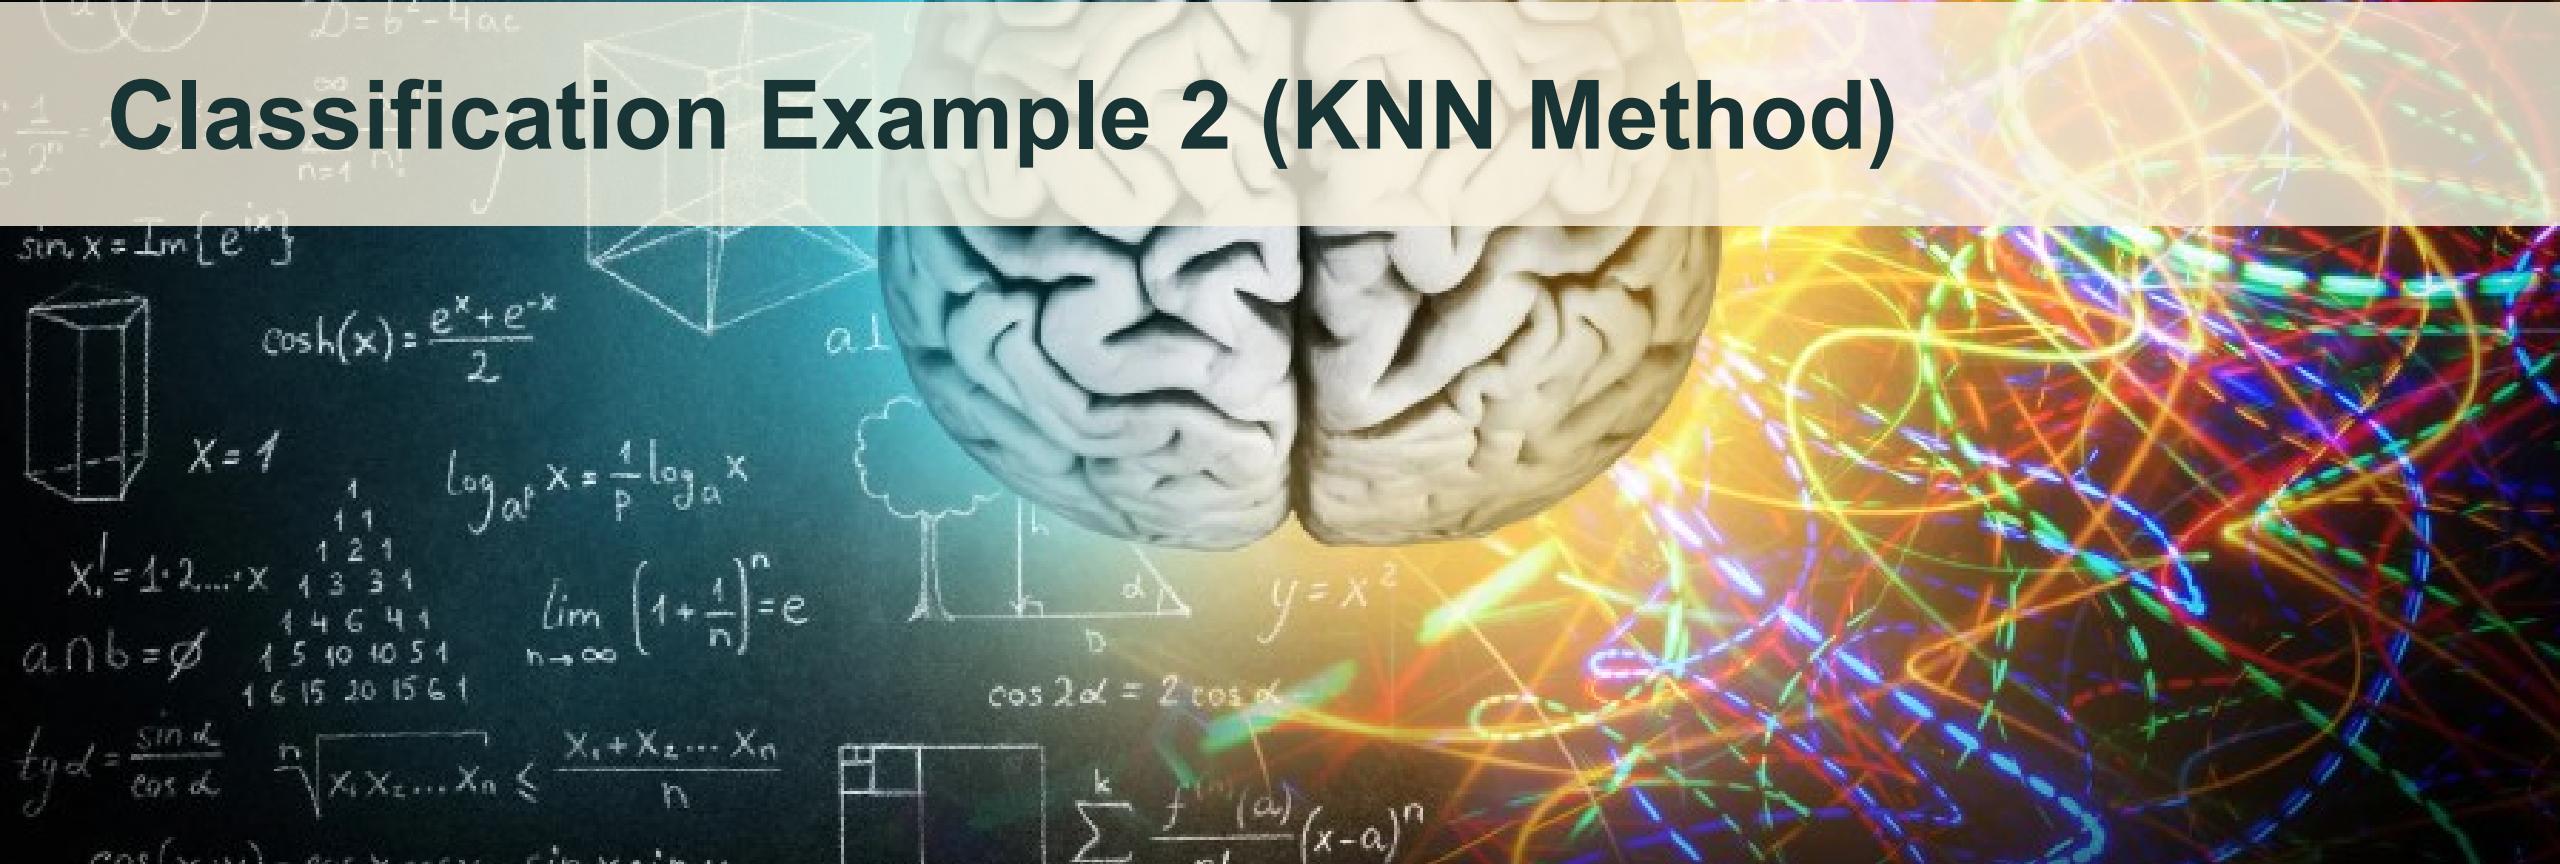
- ROC AUC = measure of classifier performance
- Perfect model: AUC = 1.0
- Random model: AUC = 0.5
- Higher AUC → better classifier overall

```
>>> from sklearn.metrics import roc_auc_score  
>>> roc_auc_score(y_train_5, y_scores)  
0.9604938554008616
```



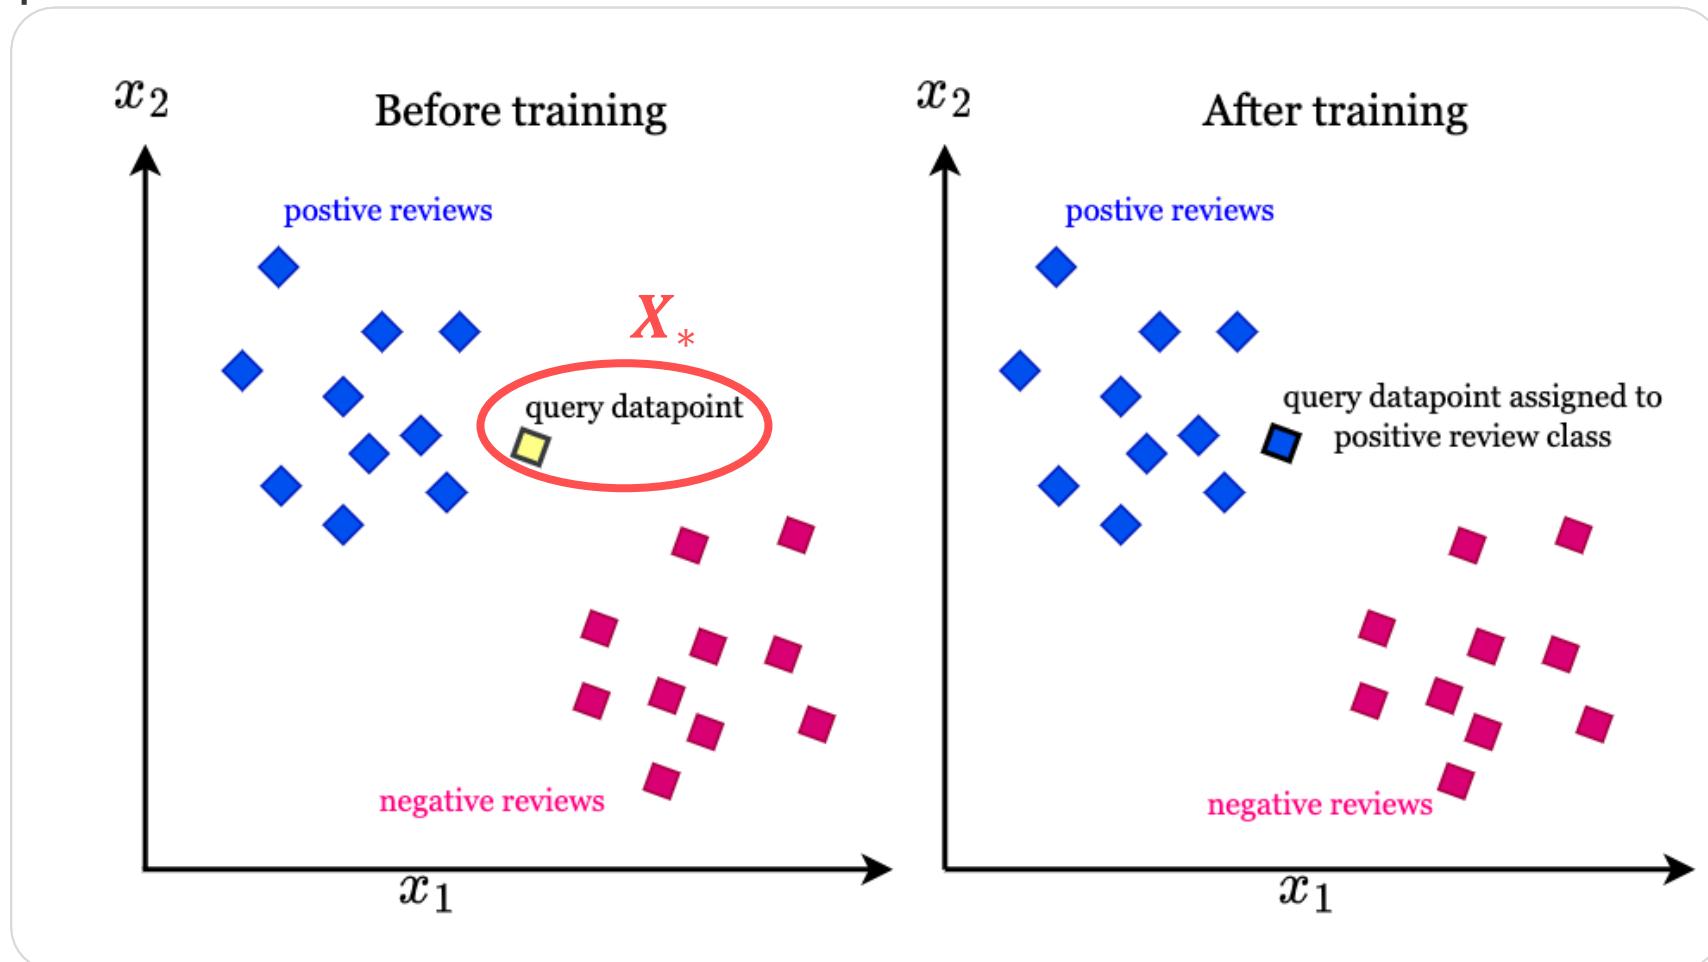


Classification Example 2 (KNN Method)



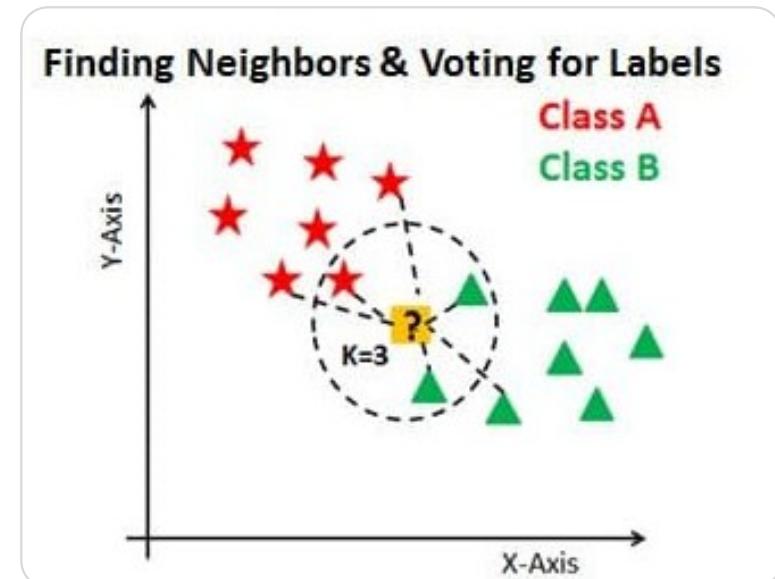
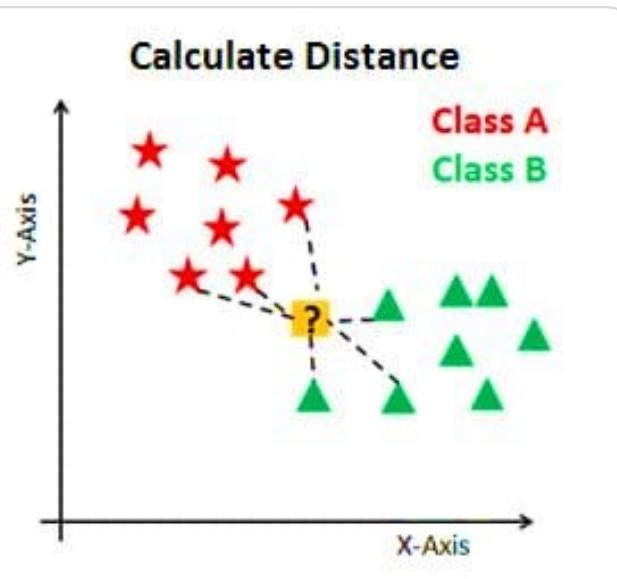
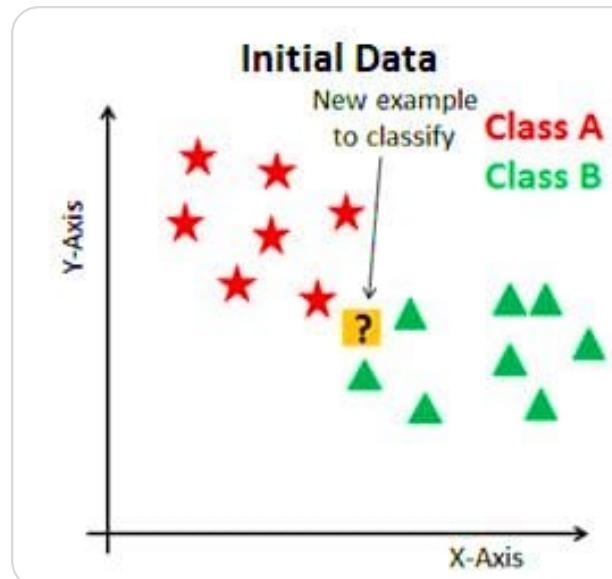
k-Nearest Neighbors (k-NN)

- Predict based on nearby k data points (k-Nearest Neighbors).
 - If a new item is similar to old ones, its output is likely similar too.
 - Non-parametric model



Working of k-NN

- Measure distance between new point and all old points
- Pick the closest k points (neighbors).
- **For regression:** take the average output
 - A new house price is predicted using prices of 3 nearby houses
- **For classification:** choose the most common output:
 - A fruit is labeled as “apple” if most nearby fruits are apples



Steps of k-NN

- Training data: (X, y) and new data: x_\star
- **Goal:** Predict the output (\hat{y}_\star) for x_\star
- **Step-1:** Measure Distance (use Euclidean distance or another method).
- **Step-2:** Pick the k closest data points (neighbors)
- **Step-3:** For regression take the average output and for classification use majority vote.

Data: Training data $\{\mathbf{x}_i, y_i\}_{i=1}^n$ and test input \mathbf{x}_\star

Result: Predicted test output $\hat{y}(\mathbf{x}_\star)$

- 1 Compute the distances $\|\mathbf{x}_i - \mathbf{x}_\star\|_2$ for all training data points $i = 1, \dots, n$
- 2 Let $N_\star = \{i : \mathbf{x}_i \text{ is one of the } k \text{ data points closest to } \mathbf{x}_\star\}$
- 3 Compute the prediction $\hat{y}(\mathbf{x}_\star)$ as

$$\hat{y}(\mathbf{x}_\star) = \begin{cases} \text{Average}\{y_j : j \in N_\star\} & \text{(Regression problems)} \\ \text{MajorityVote}\{y_j : j \in N_\star\} & \text{(Classification problems)} \end{cases}$$

Example: Predicting Colors with k-NN

- Training Data (6 Points) and **Test Data Point:** $x_\star = [1, 2]$

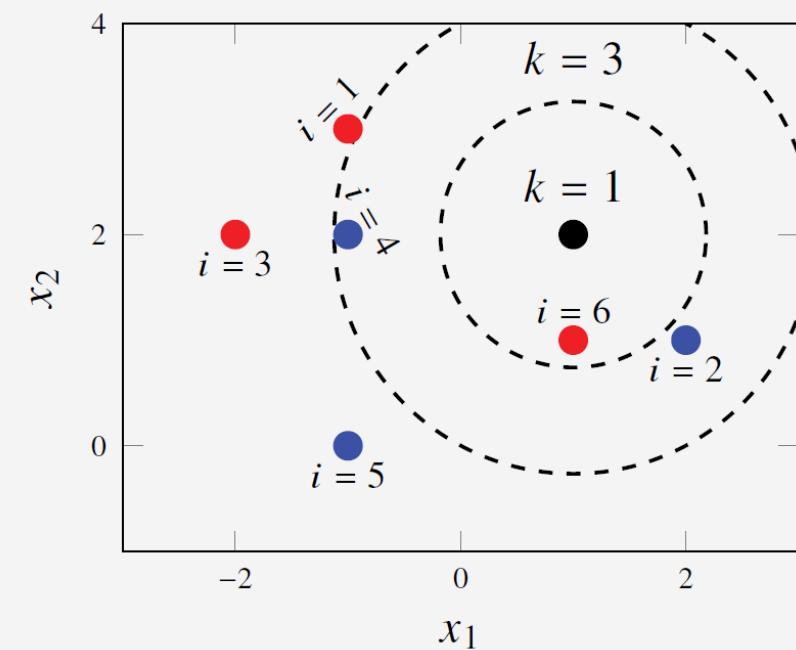
Training Data			
i	x_1	x_2	y
1	-1	3	Red
2	2	1	Blue
3	-2	2	Red
4	-1	2	Blue
5	-1	0	Blue
6	1	1	Red

Find how far each training point is from $x_\star = [1, 2]$

Calculate Distances (Euclidean)

i	$\ \mathbf{x}_i - \mathbf{x}_\star\ _2$	y_i
6	$\sqrt{1}$	Red
2	$\sqrt{2}$	Blue
4	$\sqrt{4}$	Blue
1	$\sqrt{5}$	Red
5	$\sqrt{8}$	Blue
3	$\sqrt{9}$	Red

Fig.
2.3



Prediction with $k = 1$

- Closest data point: Point 6 → Red
 - Prediction = Red

Find how far each training point is from $\mathbf{x}_\star = [1, 2]$

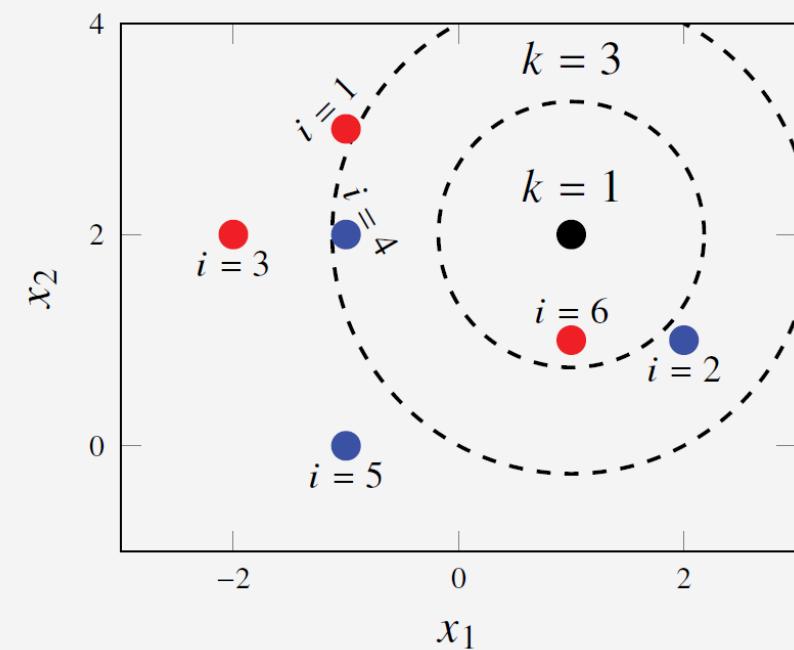
Training Data

i	x_1	x_2	y
1	-1	3	Red
2	2	1	Blue
3	-2	2	Red
4	-1	2	Blue
5	-1	0	Blue
6	1	1	Red

Calculate Distances
(Euclidean)

i	$\ \mathbf{x}_i - \mathbf{x}_\star\ _2$	y_i
6	$\sqrt{1}$	Red
2	$\sqrt{2}$	Blue
4	$\sqrt{4}$	Blue
1	$\sqrt{5}$	Red
5	$\sqrt{8}$	Blue
3	$\sqrt{9}$	Red

Fig.
2.3



Prediction with $k = 3$

- Closest data points:
 - Point 6 → Red
 - Point 2 → Blue
 - Point 4 → Blue
- Majority Vote: Blue (2), Red (1)
 - Prediction = Blue

Training Data

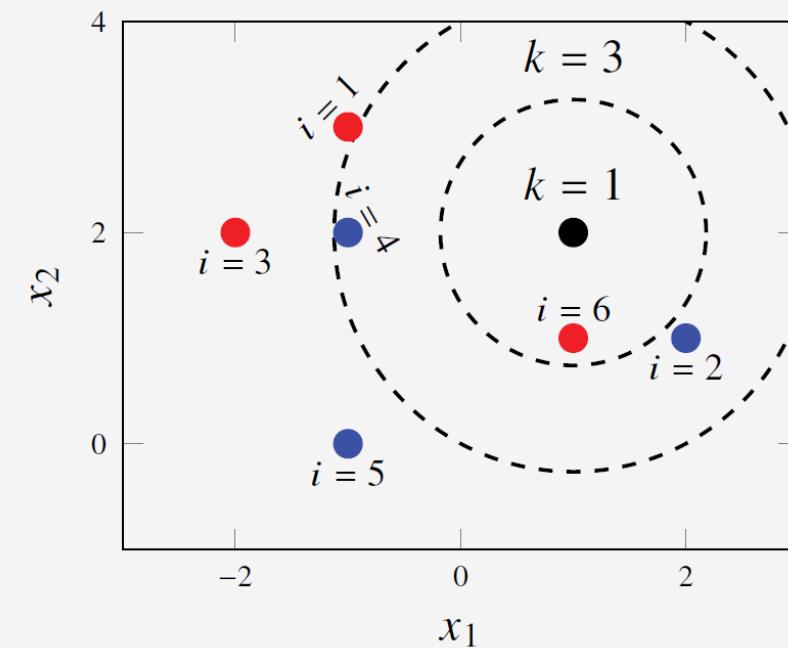
i	x_1	x_2	y
1	-1	3	Red
2	2	1	Blue
3	-2	2	Red
4	-1	2	Blue
5	-1	0	Blue
6	1	1	Red

Find how far each training point is from $\mathbf{x}_\star = [1, 2]$

Calculate Distances (Euclidean)

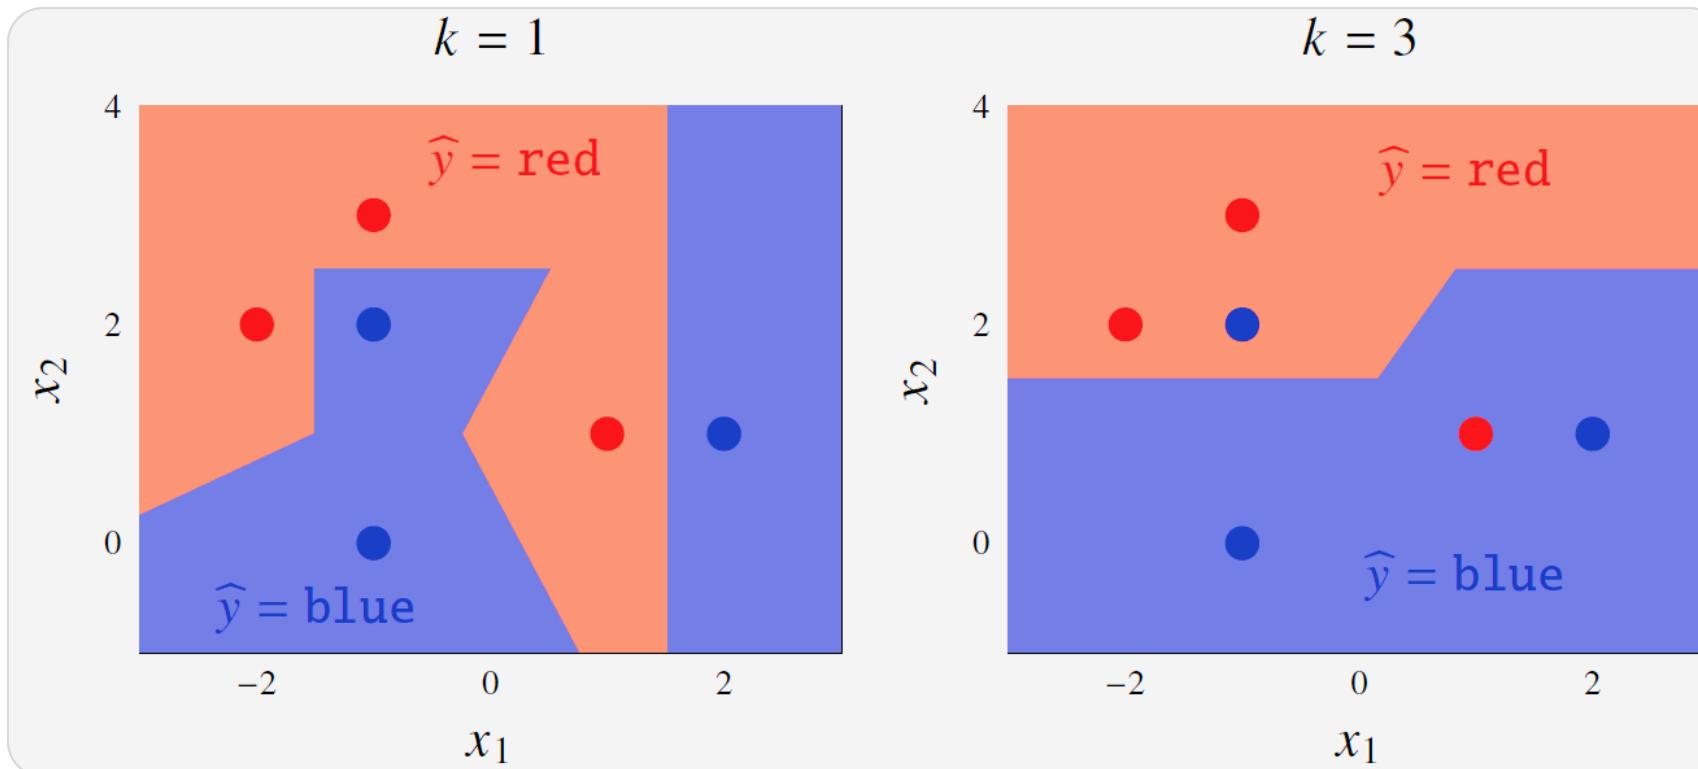
i	$\ \mathbf{x}_i - \mathbf{x}_\star\ _2$	y_i
6	$\sqrt{1}$	Red
2	$\sqrt{2}$	Blue
4	$\sqrt{4}$	Blue
1	$\sqrt{5}$	Red
5	$\sqrt{8}$	Blue
3	$\sqrt{9}$	Red

Fig.
2.3



Decision Boundaries for a Classifier

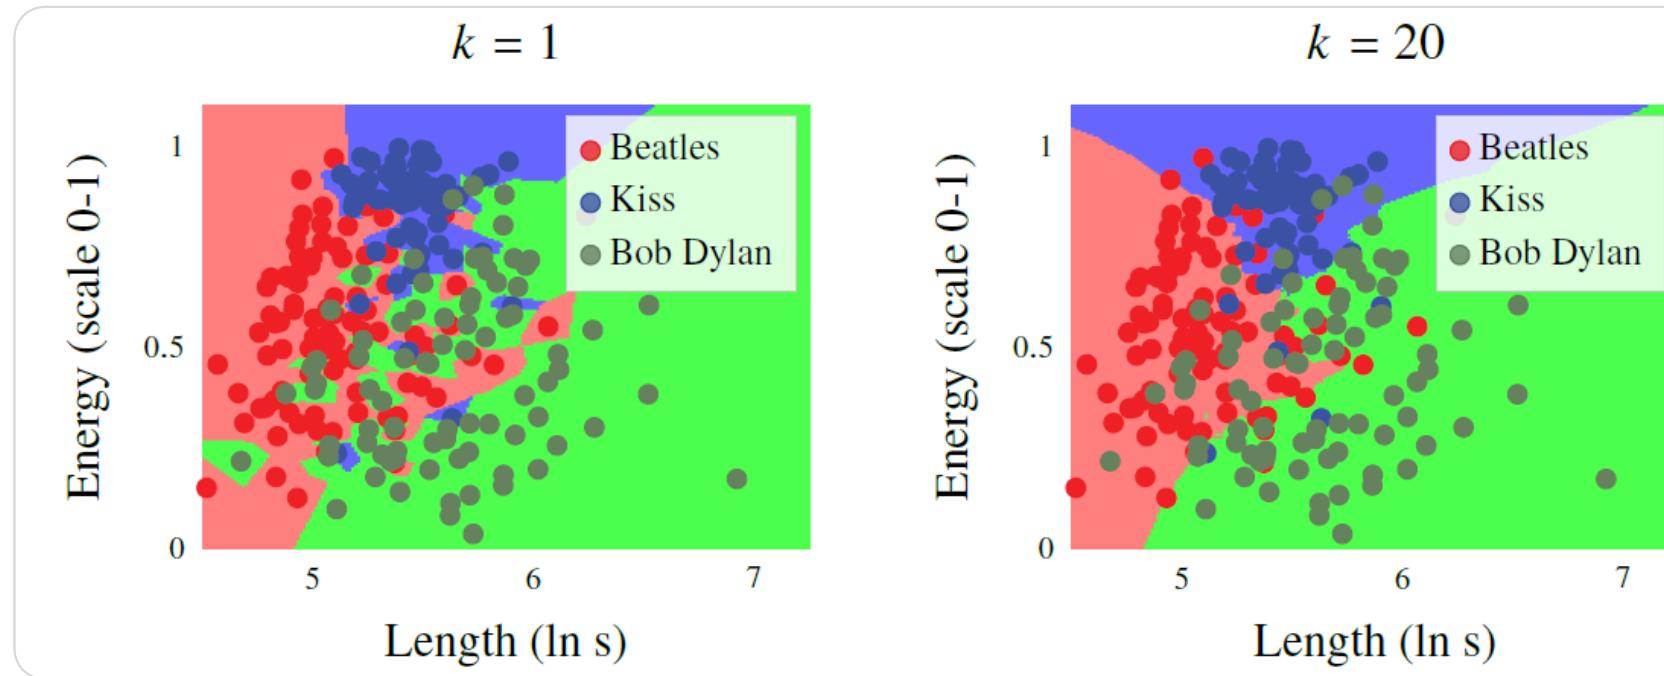
- A decision boundary is the line (or curve) where the predicted class changes.
- It separates regions where the model predicts Red or Blue.
 - For $k = 1 \rightarrow$ boundaries are more jagged
 - For $k = 3 \rightarrow$ smoother, more stable regions



What is a Good Value for K?

- K is a hyperparameter (chosen by user, not learned by model).
- Best k is found by testing different values

Small k (e.g., k = 1): Fits training data too closely (overfitting)



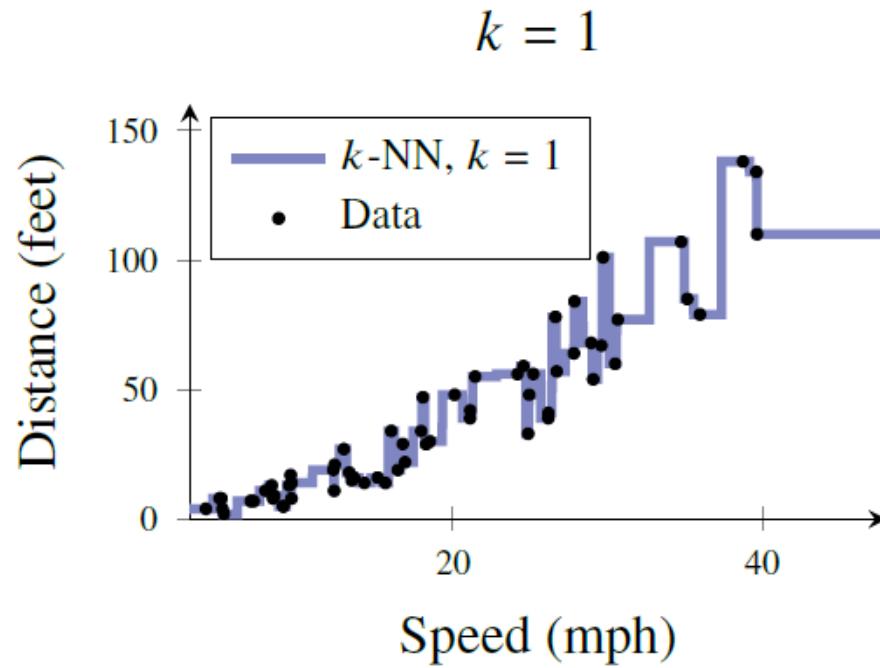
Large k (e.g., k = 20) Fits training data better for generalizing to new data

***k*-NN applied to the music classification with $K = 1$ and $K = 20$**

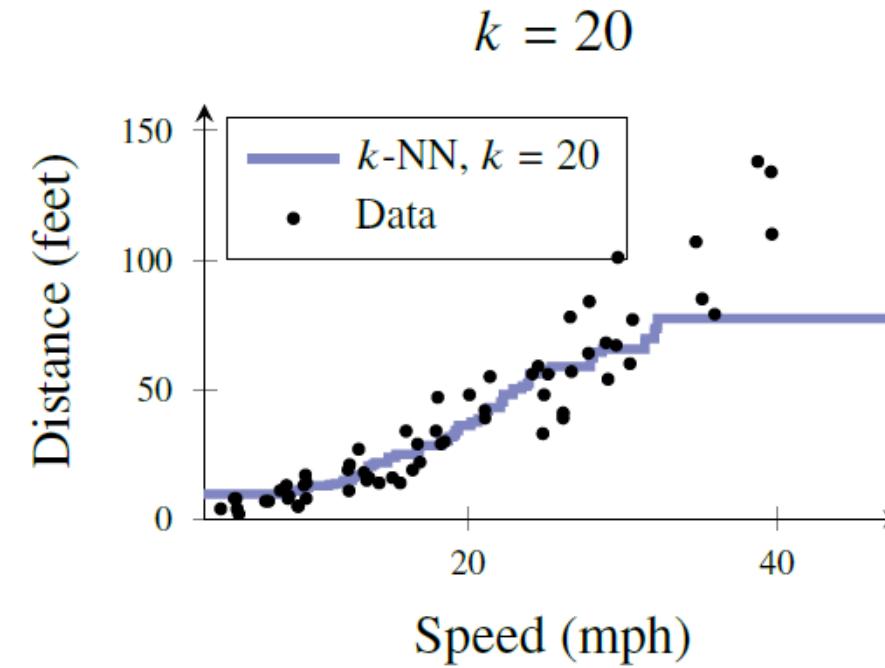
What is a Good Value for K?

- K is a hyperparameter (chosen by user, not learned by model).
- Best k is found by testing different values

Small k (e.g., k = 1): Fits training data too closely (overfitting)



Large k (e.g., k = 20) Fits training data better for generalizing to new data



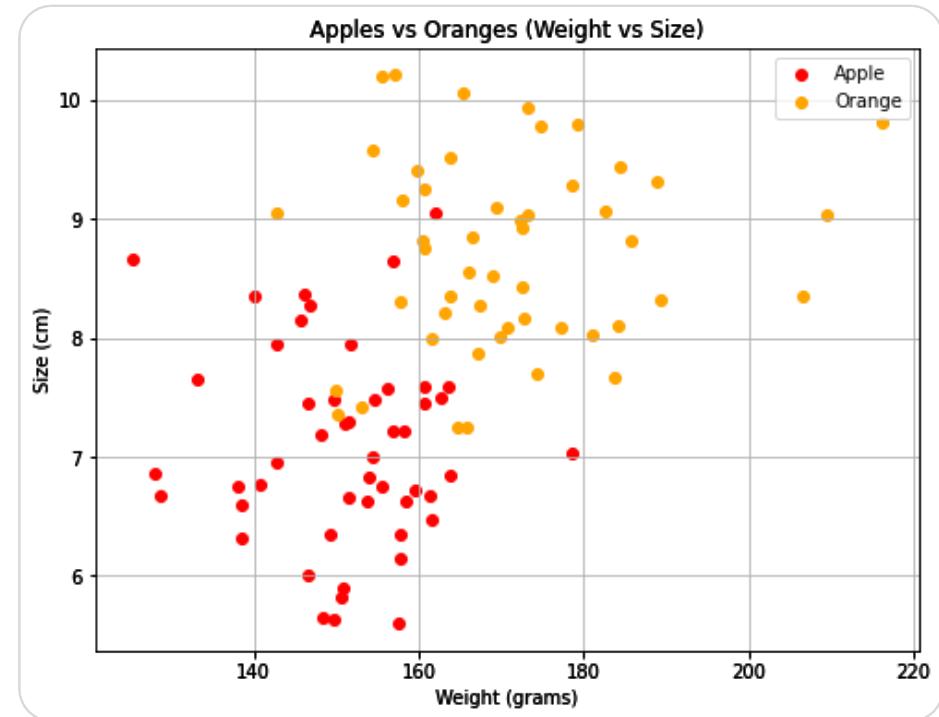
k -NN applied to the car stopping distance with $K = 1$ and $K = 20$

Input Normalization in k-NN

- k-NN uses distance to find neighbors
- If one feature has a bigger range, it will dominate the distance:
 - X1: Score [0–1]
 - X2: Height in cm [100–200] → height controls the result
- **Solution:** Normalize your input data
- Two standard ways:
 - **Min–Max scaling:** Squash into [0,1] by $f(x) = (x - \text{min}) / (\text{max-min})$
 - **Standardization:** Subtract mean and divide by standard deviation:
 $f(x) = (x - \text{mean})/\text{sdev}$

KNN in Action — Real Example

- We classify fruits   based on:
 - Weight (grams)
 - Size (diameter in cm)
- **Two classes:**
 - Apple → 0
 - Orange → 1
- **Goal:** Predict fruit type using KNN



Create simple real-world-like data for two fruits

```
# =====
# 2 Create simple real-world-like data for two fruits
# =====
# Feature 1: Weight (grams)
# Feature 2: Size (diameter in cm)

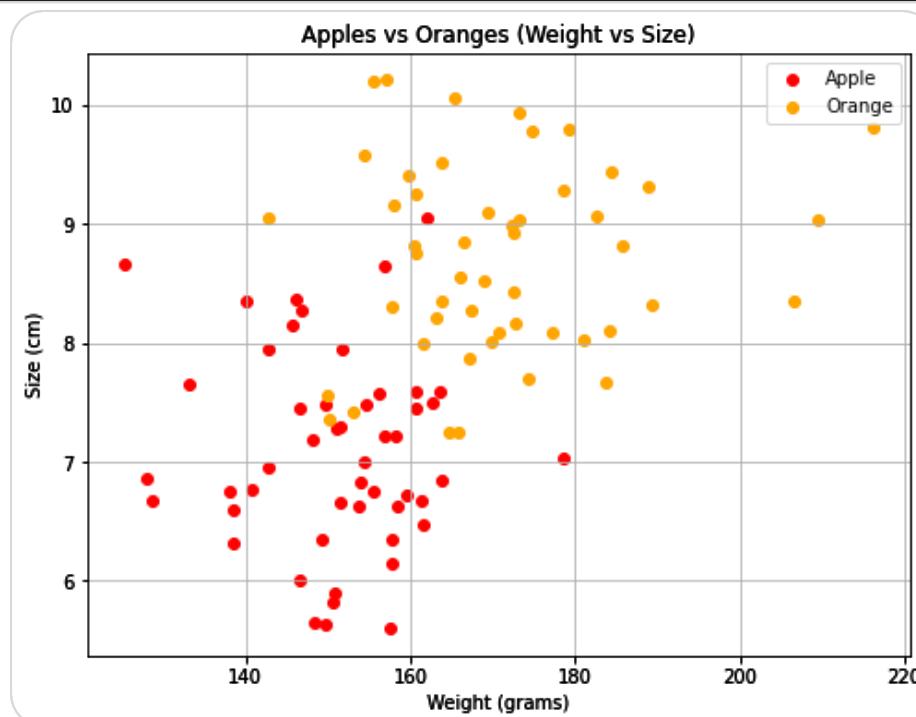
# Class 0 → Apples: lighter and smaller
apple_weight = np.random.normal(150, 10, 50)      # mean=150g, std=10
apple_size   = np.random.normal(7, 0.8, 50)        # mean=7cm, std=0.8

# Class 1 → Oranges: heavier and larger
orange_weight = np.random.normal(200, 15, 50)      # mean=200g, std=15
orange_size   = np.random.normal(8.5, 0.8, 50)       # mean=8.5cm, std=0.8

# Combine both classes into one dataset
X = np.column_stack((
    np.concatenate([apple_weight, orange_weight]),
    np.concatenate([apple_size, orange_size])
))
y = np.concatenate([np.zeros(50), np.ones(50)])      # 0 = Apple, 1 = Orange
```

Visualize Data

```
# Visualize the data
plt.figure(figsize=(8,6))
plt.scatter(apple_weight, apple_size, color='red', label='Apple')
plt.scatter(orange_weight, orange_size, color='orange', label='Orange')
plt.title("Apples vs Oranges (Weight vs Size)")
plt.xlabel("Weight (grams)")
plt.ylabel("Size (cm)")
plt.legend()
plt.grid(True)
plt.show()
```



Train KNN Model

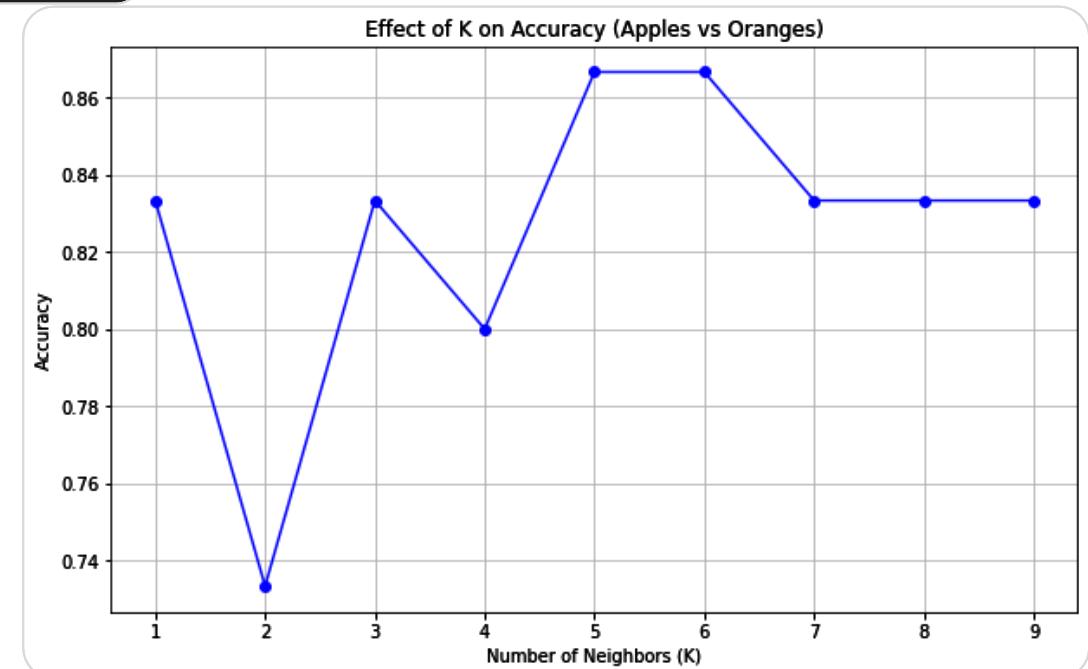
```
# =====
# 3 Split data into training and test sets
# =====
X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                    test_size=0.3,
                                                    random_state=42)
# =====
# 5 Check accuracy for different K values
# =====
k_values = []
accuracies = []

for k in range(1, 10):
    model = KNeighborsClassifier(n_neighbors=k)
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    acc = accuracy_score(y_test, y_pred)

    k_values.append(k)
    accuracies.append(acc)
    print(f"K = {k} --> Accuracy = {acc:.2f}")
```

Check Accuracy for Different K

```
# =====
# 6 Visualize how Accuracy changes with K
# =====
plt.figure(figsize=(10,6))
plt.plot(k_values, accuracies, marker='o', linestyle='-', color='blue')
plt.title("Effect of K on Accuracy (Apples vs Oranges)")
plt.xlabel("Number of Neighbors (K)")
plt.ylabel("Accuracy")
plt.grid(True)
plt.show()
```



Visualizing Decision Boundary

```
# =====
# 4 Define function to plot decision boundary
# =====
def plot_knn(k):
    model = KNeighborsClassifier(n_neighbors=k)
    model.fit(X_train, y_train)

    # Define limits of the plot
    x_min, x_max = X[:, 0].min() - 10, X[:, 0].max() + 10
    y_min, y_max = X[:, 1].min() - 0.5, X[:, 1].max() + 0.5

    # Create a mesh grid across the feature space
    xx, yy = np.meshgrid(np.linspace(x_min, x_max, 200),
                         np.linspace(y_min, y_max, 200))

    # Combine grid points into a single array for prediction
    # np.c_ joins the flattened (raveled) coordinate arrays
    Z = model.predict(np.c_[xx.ravel(), yy.ravel()])

    # Reshape predictions back into 2D for contour plotting
    Z = Z.reshape(xx.shape)

    # Plot decision regions
    plt.figure(figsize=(6,4))
    plt.contourf(xx, yy, Z, cmap='bwr', alpha=0.2)
```

```
# Plot actual data points
plt.scatter(X_train[y_train==0, 0], X_train[y_train==0, 1],
            color='red', label='Apple (train)', edgecolors='k')
plt.scatter(X_train[y_train==1, 0], X_train[y_train==1, 1],
            color='orange', label='Orange (train)', edgecolors='k')
plt.scatter(X_test[y_test==0, 0], X_test[y_test==0, 1],
            color='pink', label='Apple (test)', edgecolors='k', marker='^')
plt.scatter(X_test[y_test==1, 0], X_test[y_test==1, 1],
            color='gold', label='Orange (test)', edgecolors='k', marker='^')

plt.title(f"KNN Decision Boundary (k = {k})")
plt.xlabel("Weight (grams)")
plt.ylabel("Size (cm)")
plt.legend()
plt.show()

# =====
# 7 Visualize decision boundaries for some K values
# =====
for k in [1, 3, 5, 10]:
    plot_knn(k)
```

Visualizing Decision Boundary

- Small $K \rightarrow$ high variance (overfitting)
- Large $K \rightarrow$ high bias (underfitting)
- Best $K \rightarrow$ balance between both
- Visual + accuracy helps find optimal K

